

# **Fundamentos de Segurança Informática (FSI)**

**2024/2025 - LEIC**

## **Software Security (Part 1)**

**Hugo Pacheco**  
[hpacheco@fc.up.pt](mailto:hpacheco@fc.up.pt)

# The Creature vs The Creator

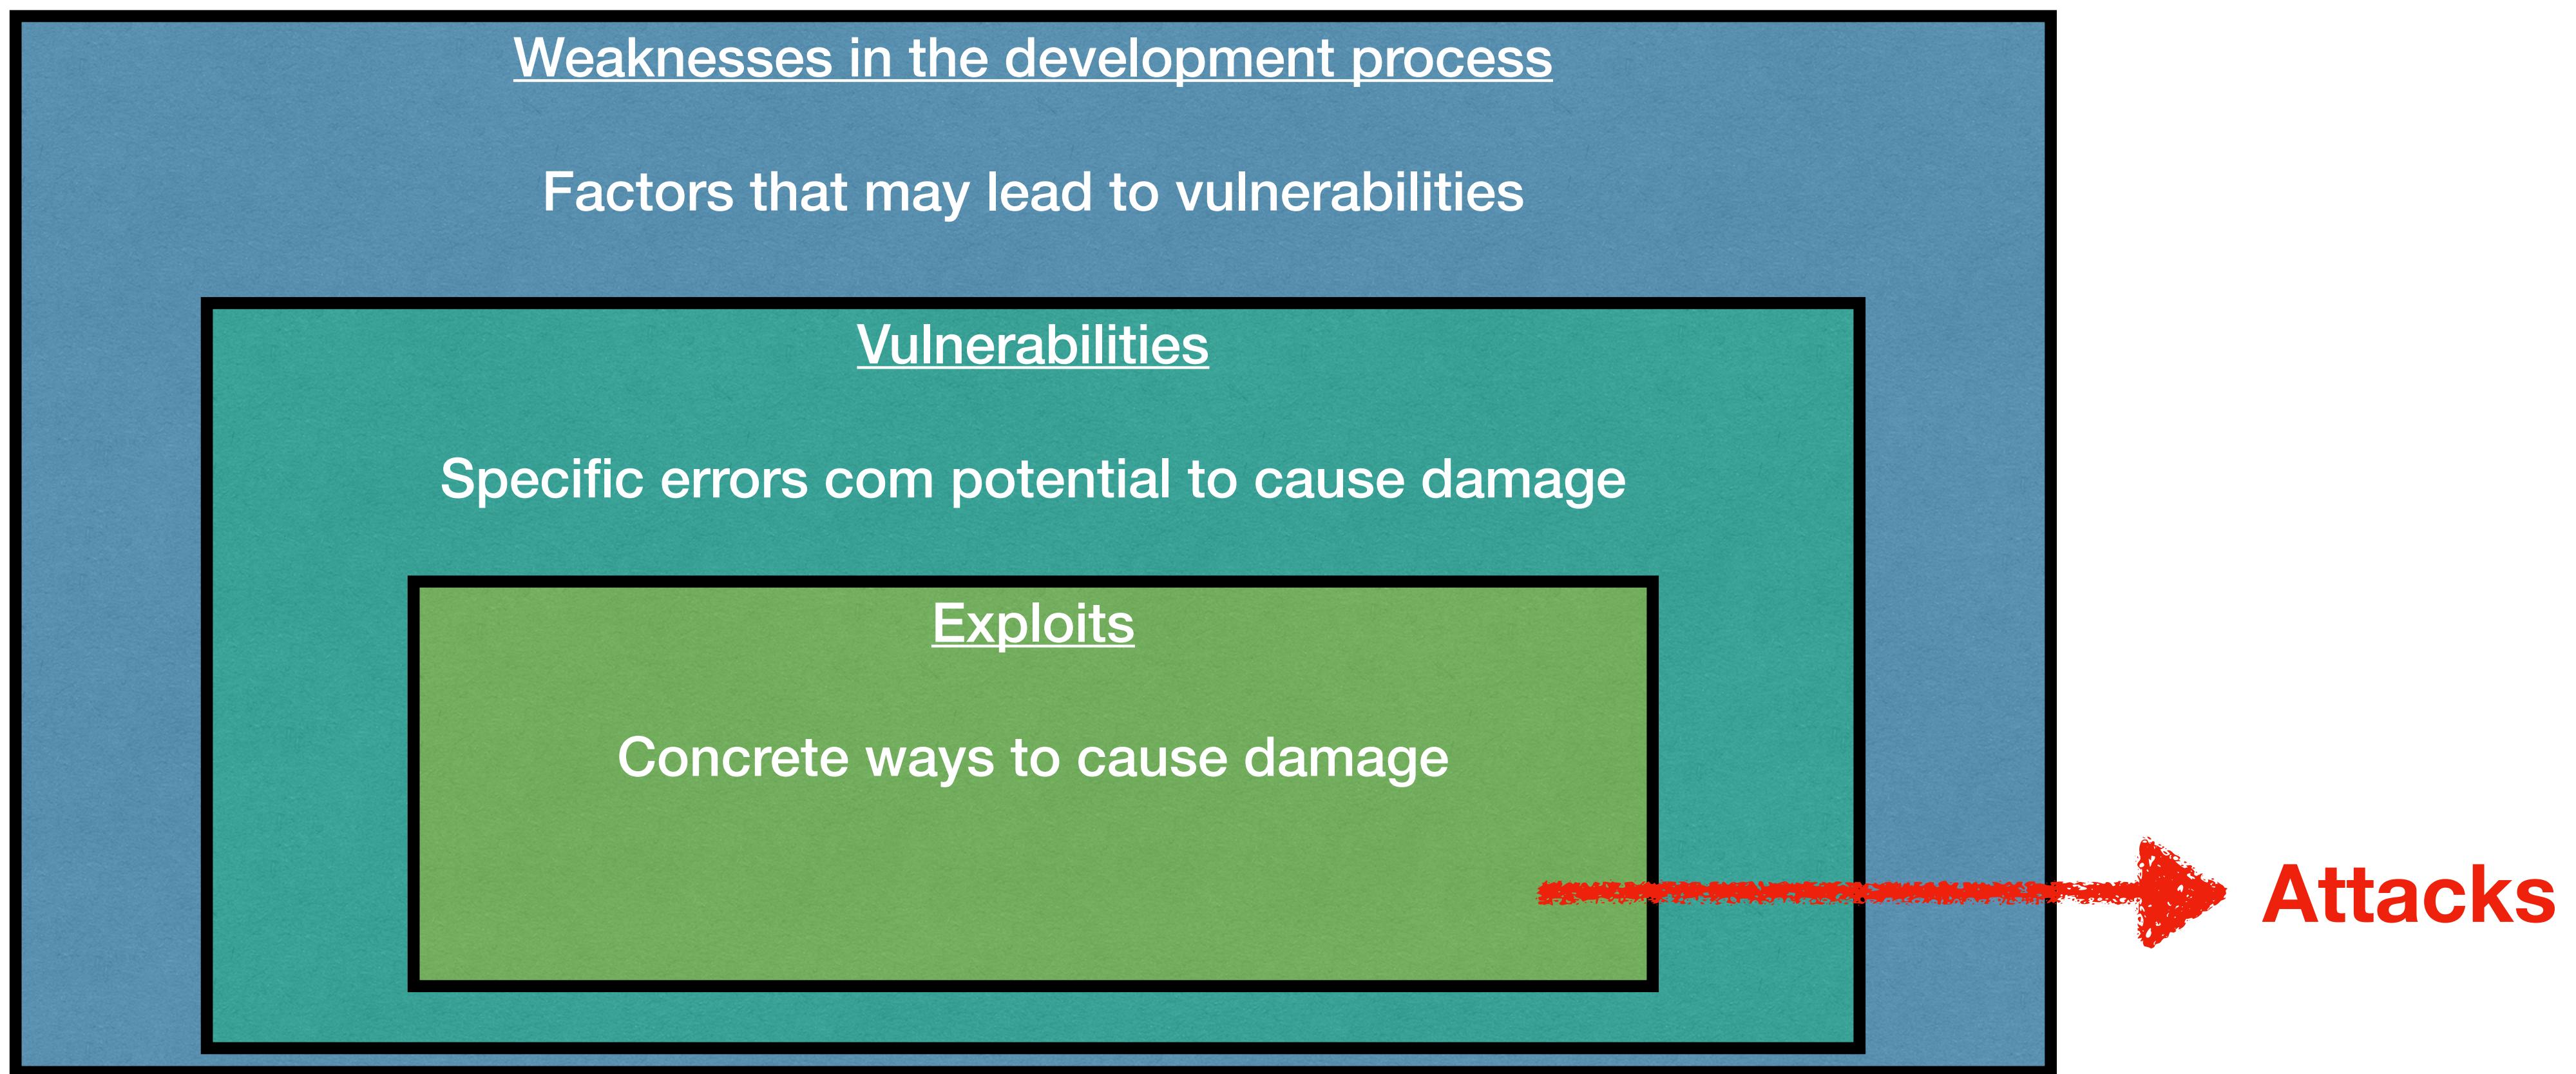
Software

Us



# The Creature vs The Creator

- Software can be manipulated to turn against us

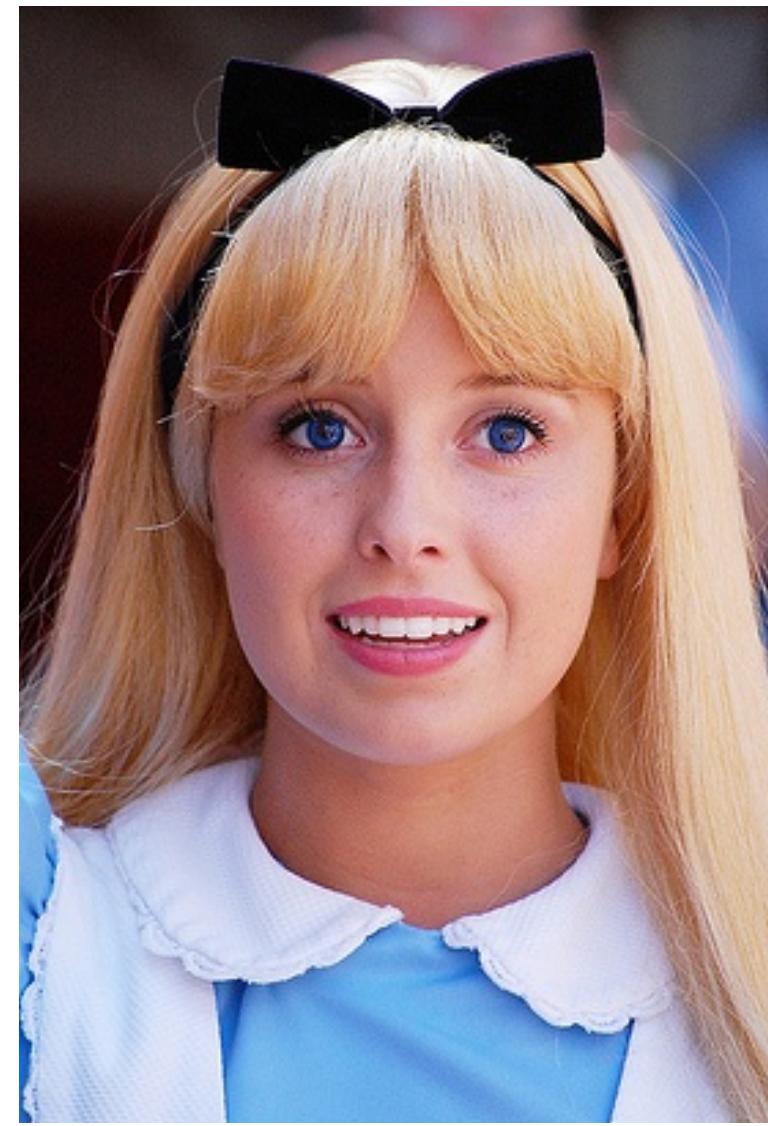


# Learning with history

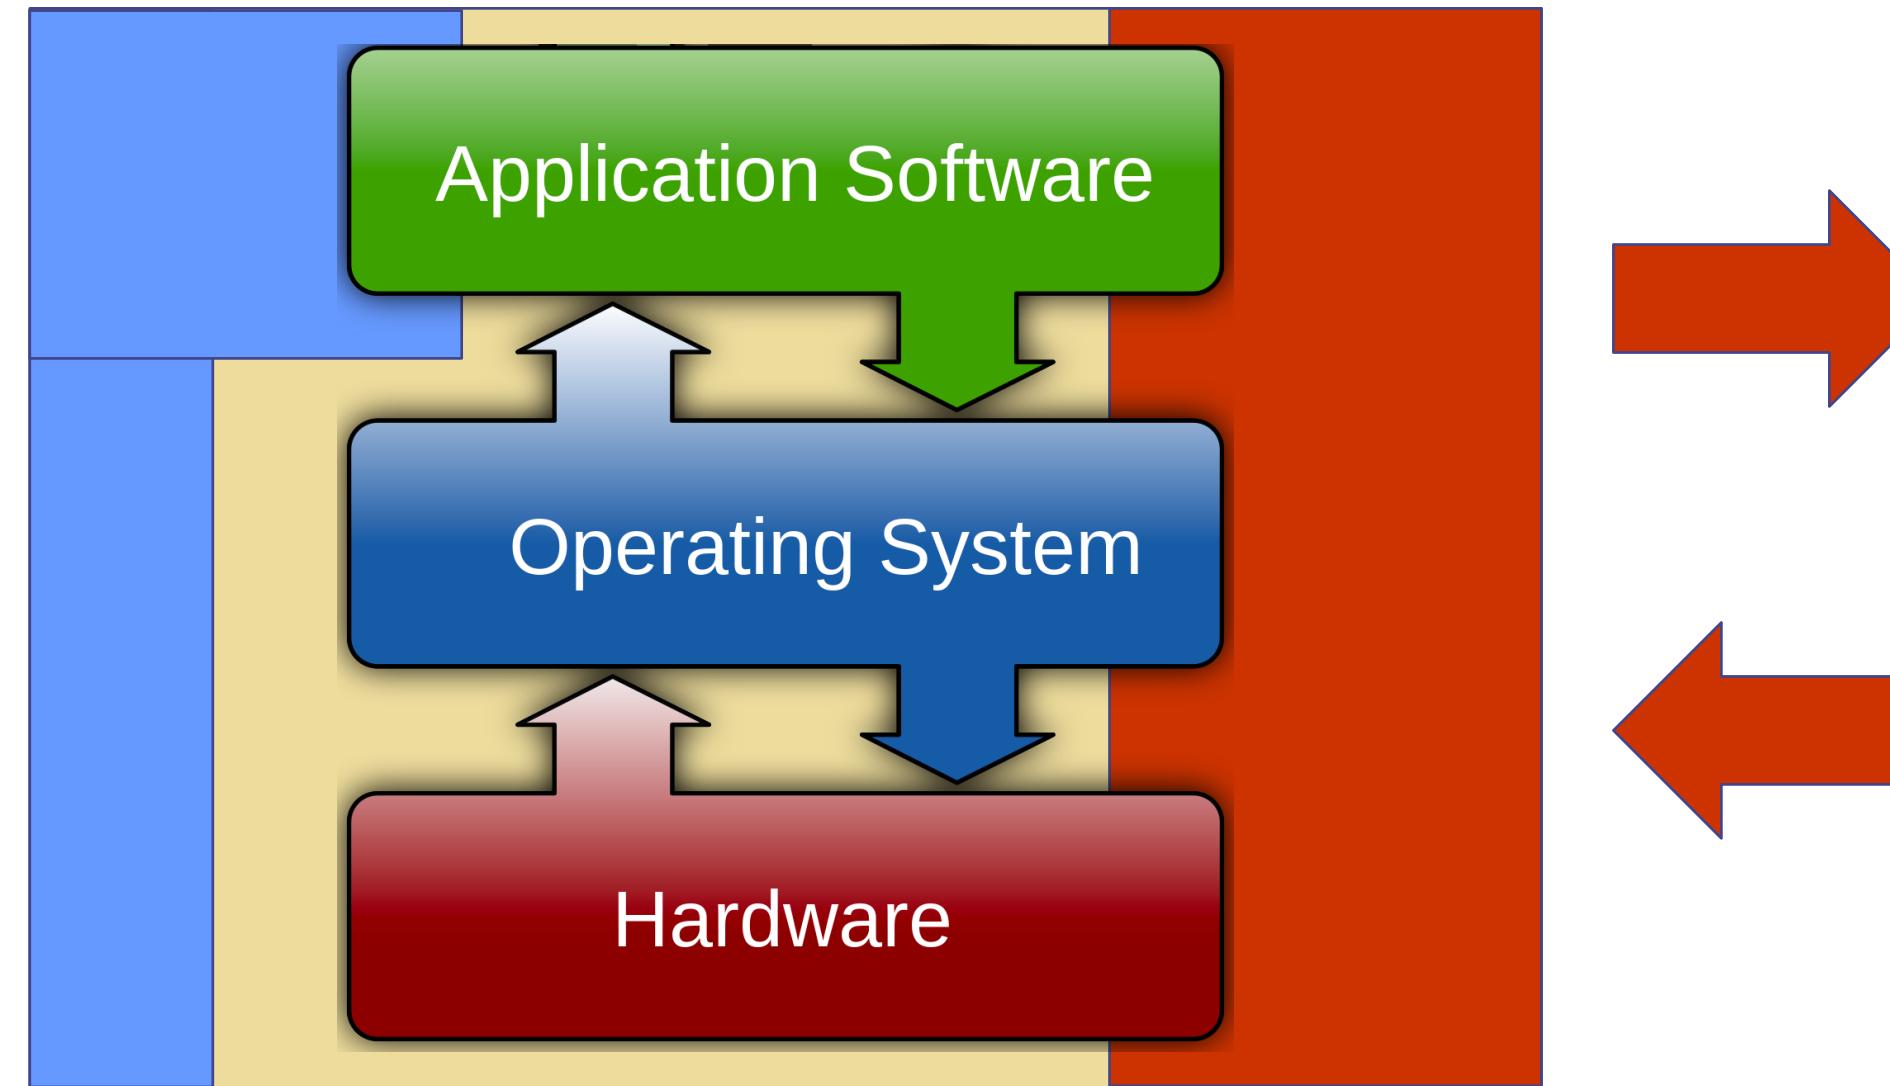
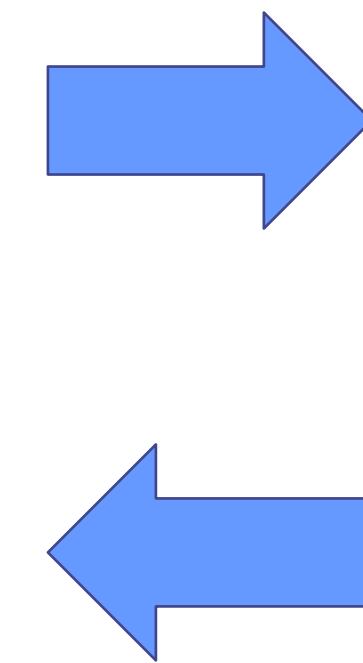
- It is important to study vulnerabilities and previous attacks
  - Identifying possible vulnerabilities and solutions
  - Incentive for fixing:
    - knowledge + flaws = negligence
  - Characterising threats:
    - Assessing the severity of vulnerabilities in a specific context
  - Help users assessing risk



# Software Security



Alice



Mallory

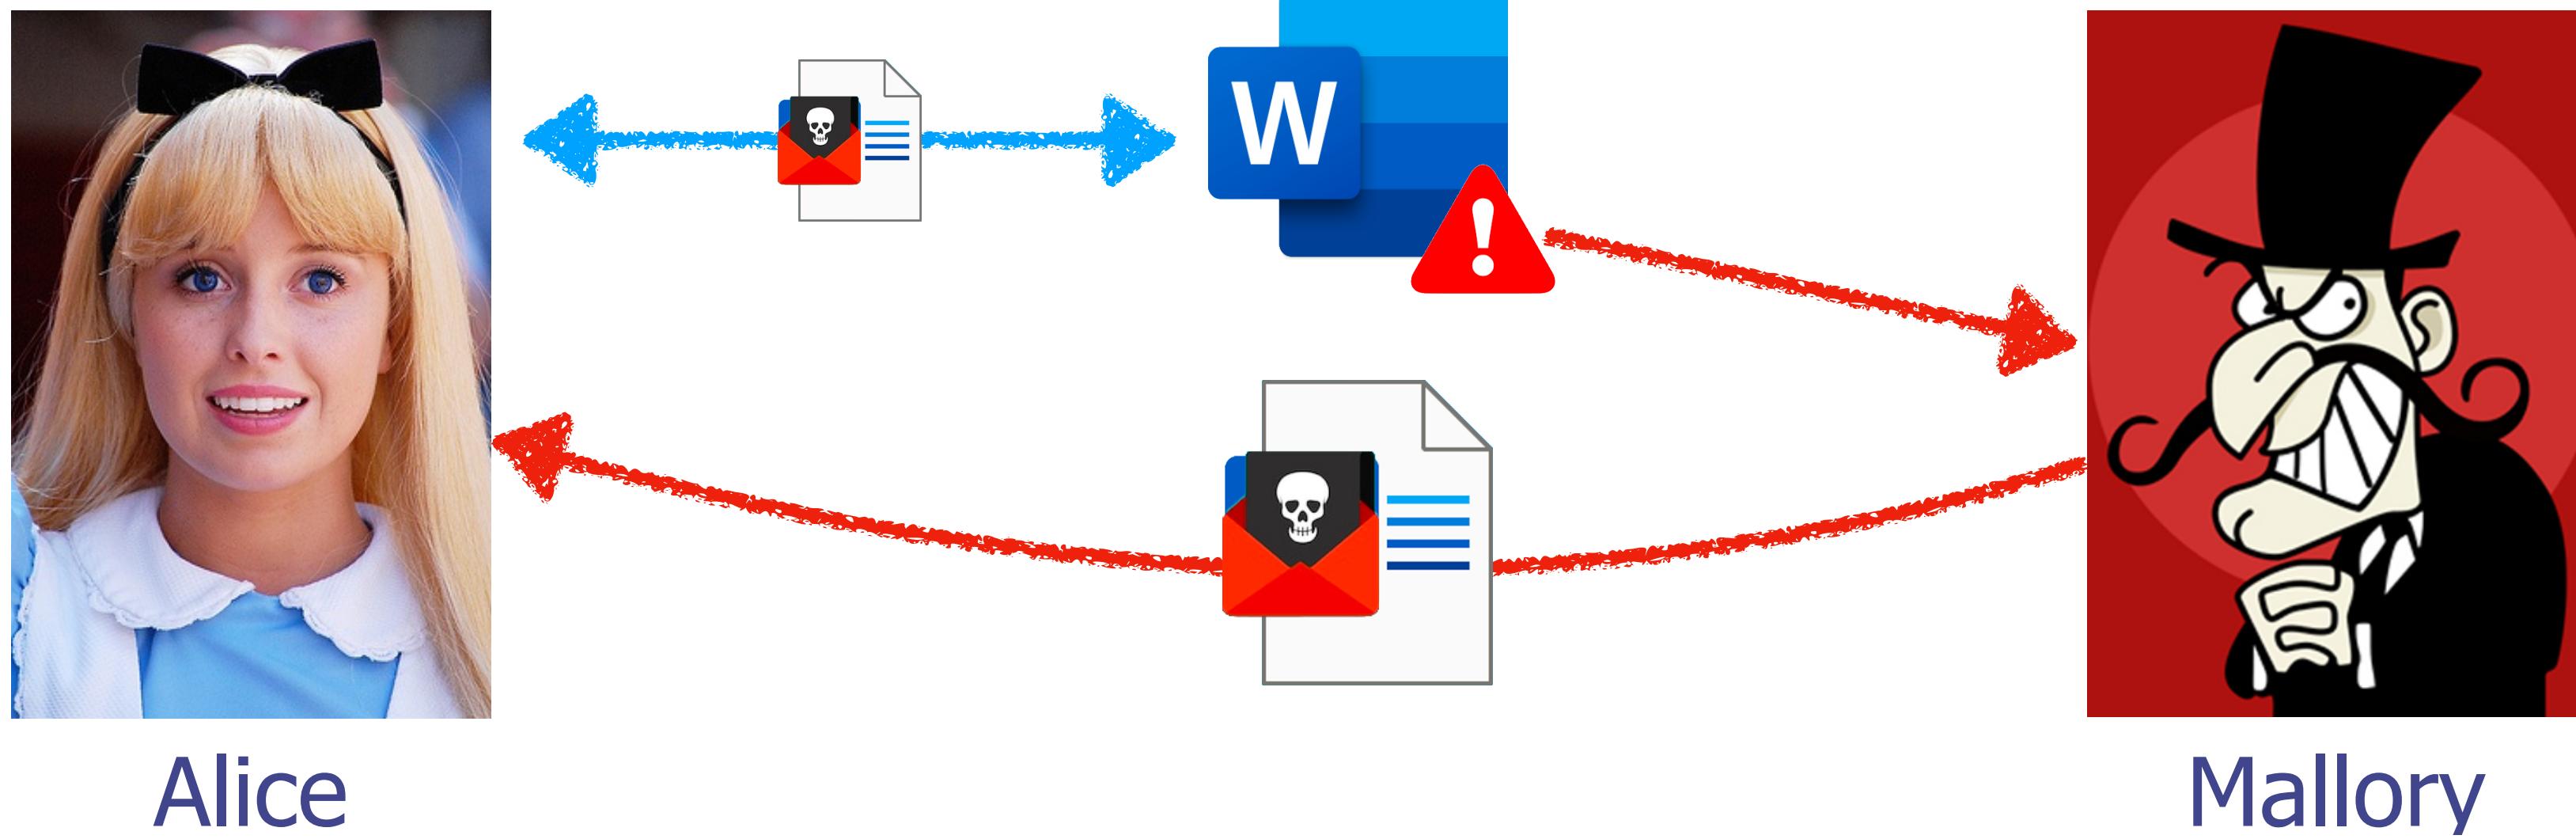
- Honest user (Alice) interacts with a software application
- Malicious attacker (Mallory) tries to exploit a software (or OS/HW) vulnerability
- Exploiting programming errors / language / compiler / execution environment

# Software Security

- Software executes in a computational platform
  - Dynamic (memory) and persistent (disk) storage
  - Operating System
  - Hardware + Interfaces with external devices
- It is always a potential weapon against that platform:
  - It may allow to **gain control** over that platform
  - Can then be used for several purposes: DoS, SpyWare, Jailbreak, etc.

# Why gaining control?

## Is crashing MS Word a security vulnerability?



- Many software vulnerabilities are violations of “control flow integrity”: **attacker alters the control flow of the program** to execute his code in **our machine**
- Threat model: **user code** receives **input from outside** the security perimeter

# The infamous “buffer overflow” & friends

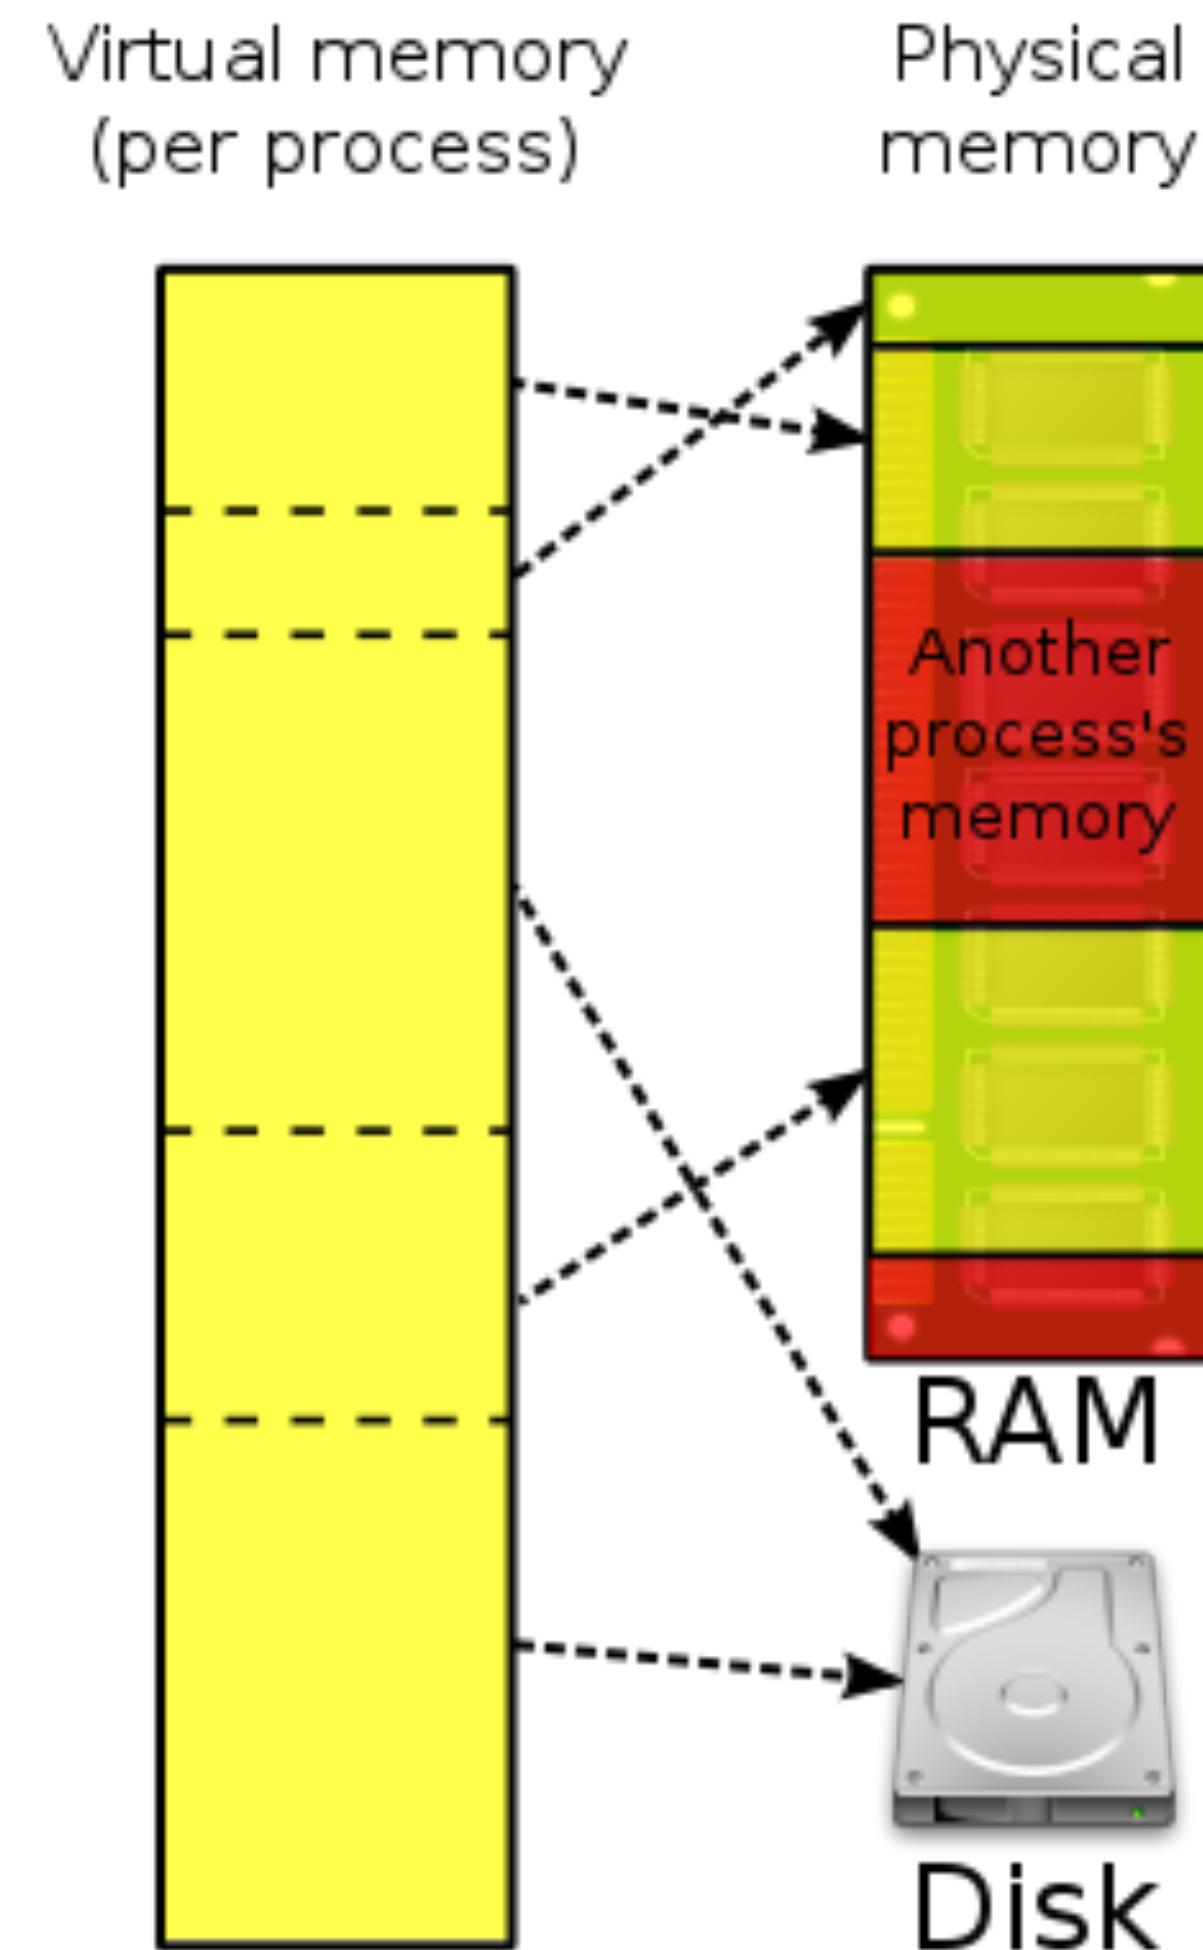


- One of the simplest (conceptually!) examples of gaining of control; belongs to a large class of memory management errors in (not only) C
- When the program “accesses outside” of its allocated memory

Rank	ID	Name	Score	CVEs in KEV	Rank Change vs. 2022
1	<a href="#">CWE-787</a>	Out-of-bounds Write	63.72	70	0
2	<a href="#">CWE-79</a>	Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')	45.54	4	0
3	<a href="#">CWE-89</a>	Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection')	34.27	6	0
4	<a href="#">CWE-416</a>	Use After Free	16.71	44	+3
5	<a href="#">CWE-78</a>	Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection')	15.65	23	+1
6	<a href="#">CWE-20</a>	Improper Input Validation	15.50	35	-2
7	<a href="#">CWE-125</a>	Out-of-bounds Read	14.60	2	-2
8	<a href="#">CWE-22</a>	Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal')	14.11	16	0
9	<a href="#">CWE-352</a>	Cross-Site Request Forgery (CSRF)	11.73	0	0
10	<a href="#">CWE-434</a>	Unrestricted Upload of File with Dangerous Type	10.41	5	0
11	<a href="#">CWE-862</a>	Missing Authorization	6.90	0	+5
12	<a href="#">CWE-476</a>	NULL Pointer Dereference	6.59	0	-1
13	<a href="#">CWE-287</a>	Improper Authentication	6.39	10	+1
14	<a href="#">CWE-190</a>	Integer Overflow or Wraparound	5.89	4	-1

# Memory Model

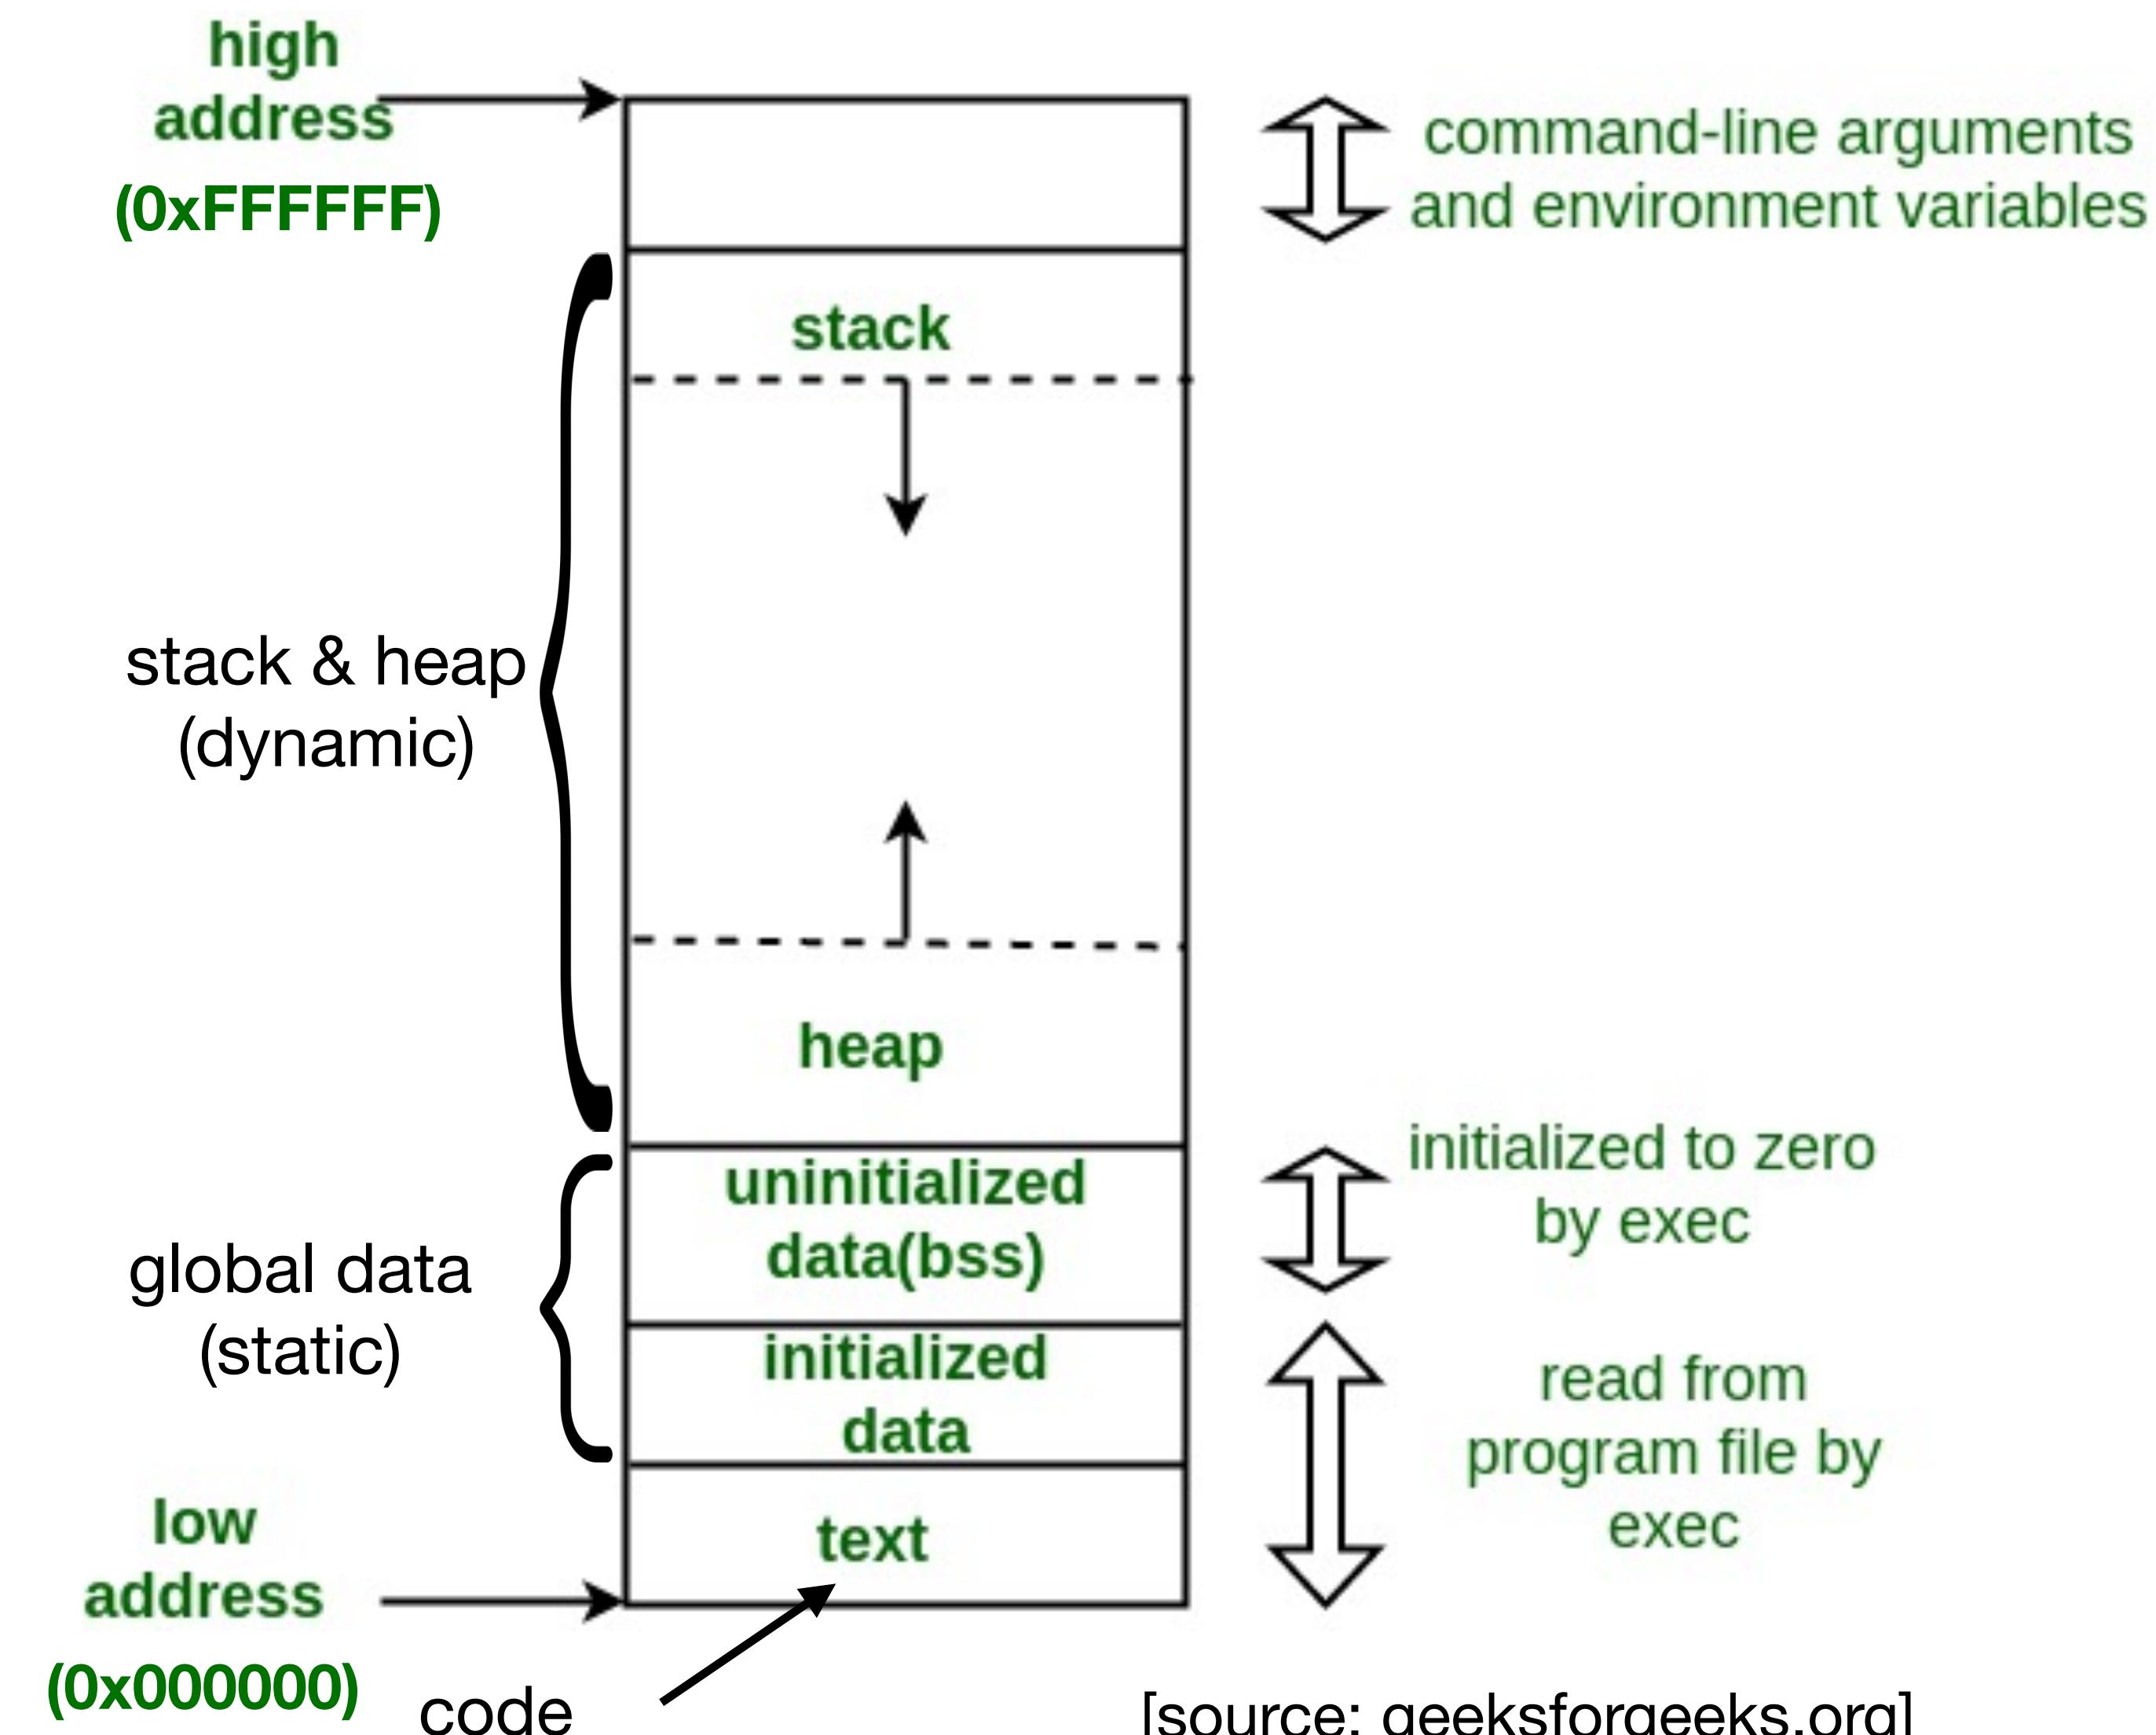
- Our examples will assume a Linux system
- In other OSes there may be small inessential differences
- Each process has access to virtual memory, that is “simulated” by the OS
- OS manages the virtual memory (e.g., paging)



[source: Wikipedia]

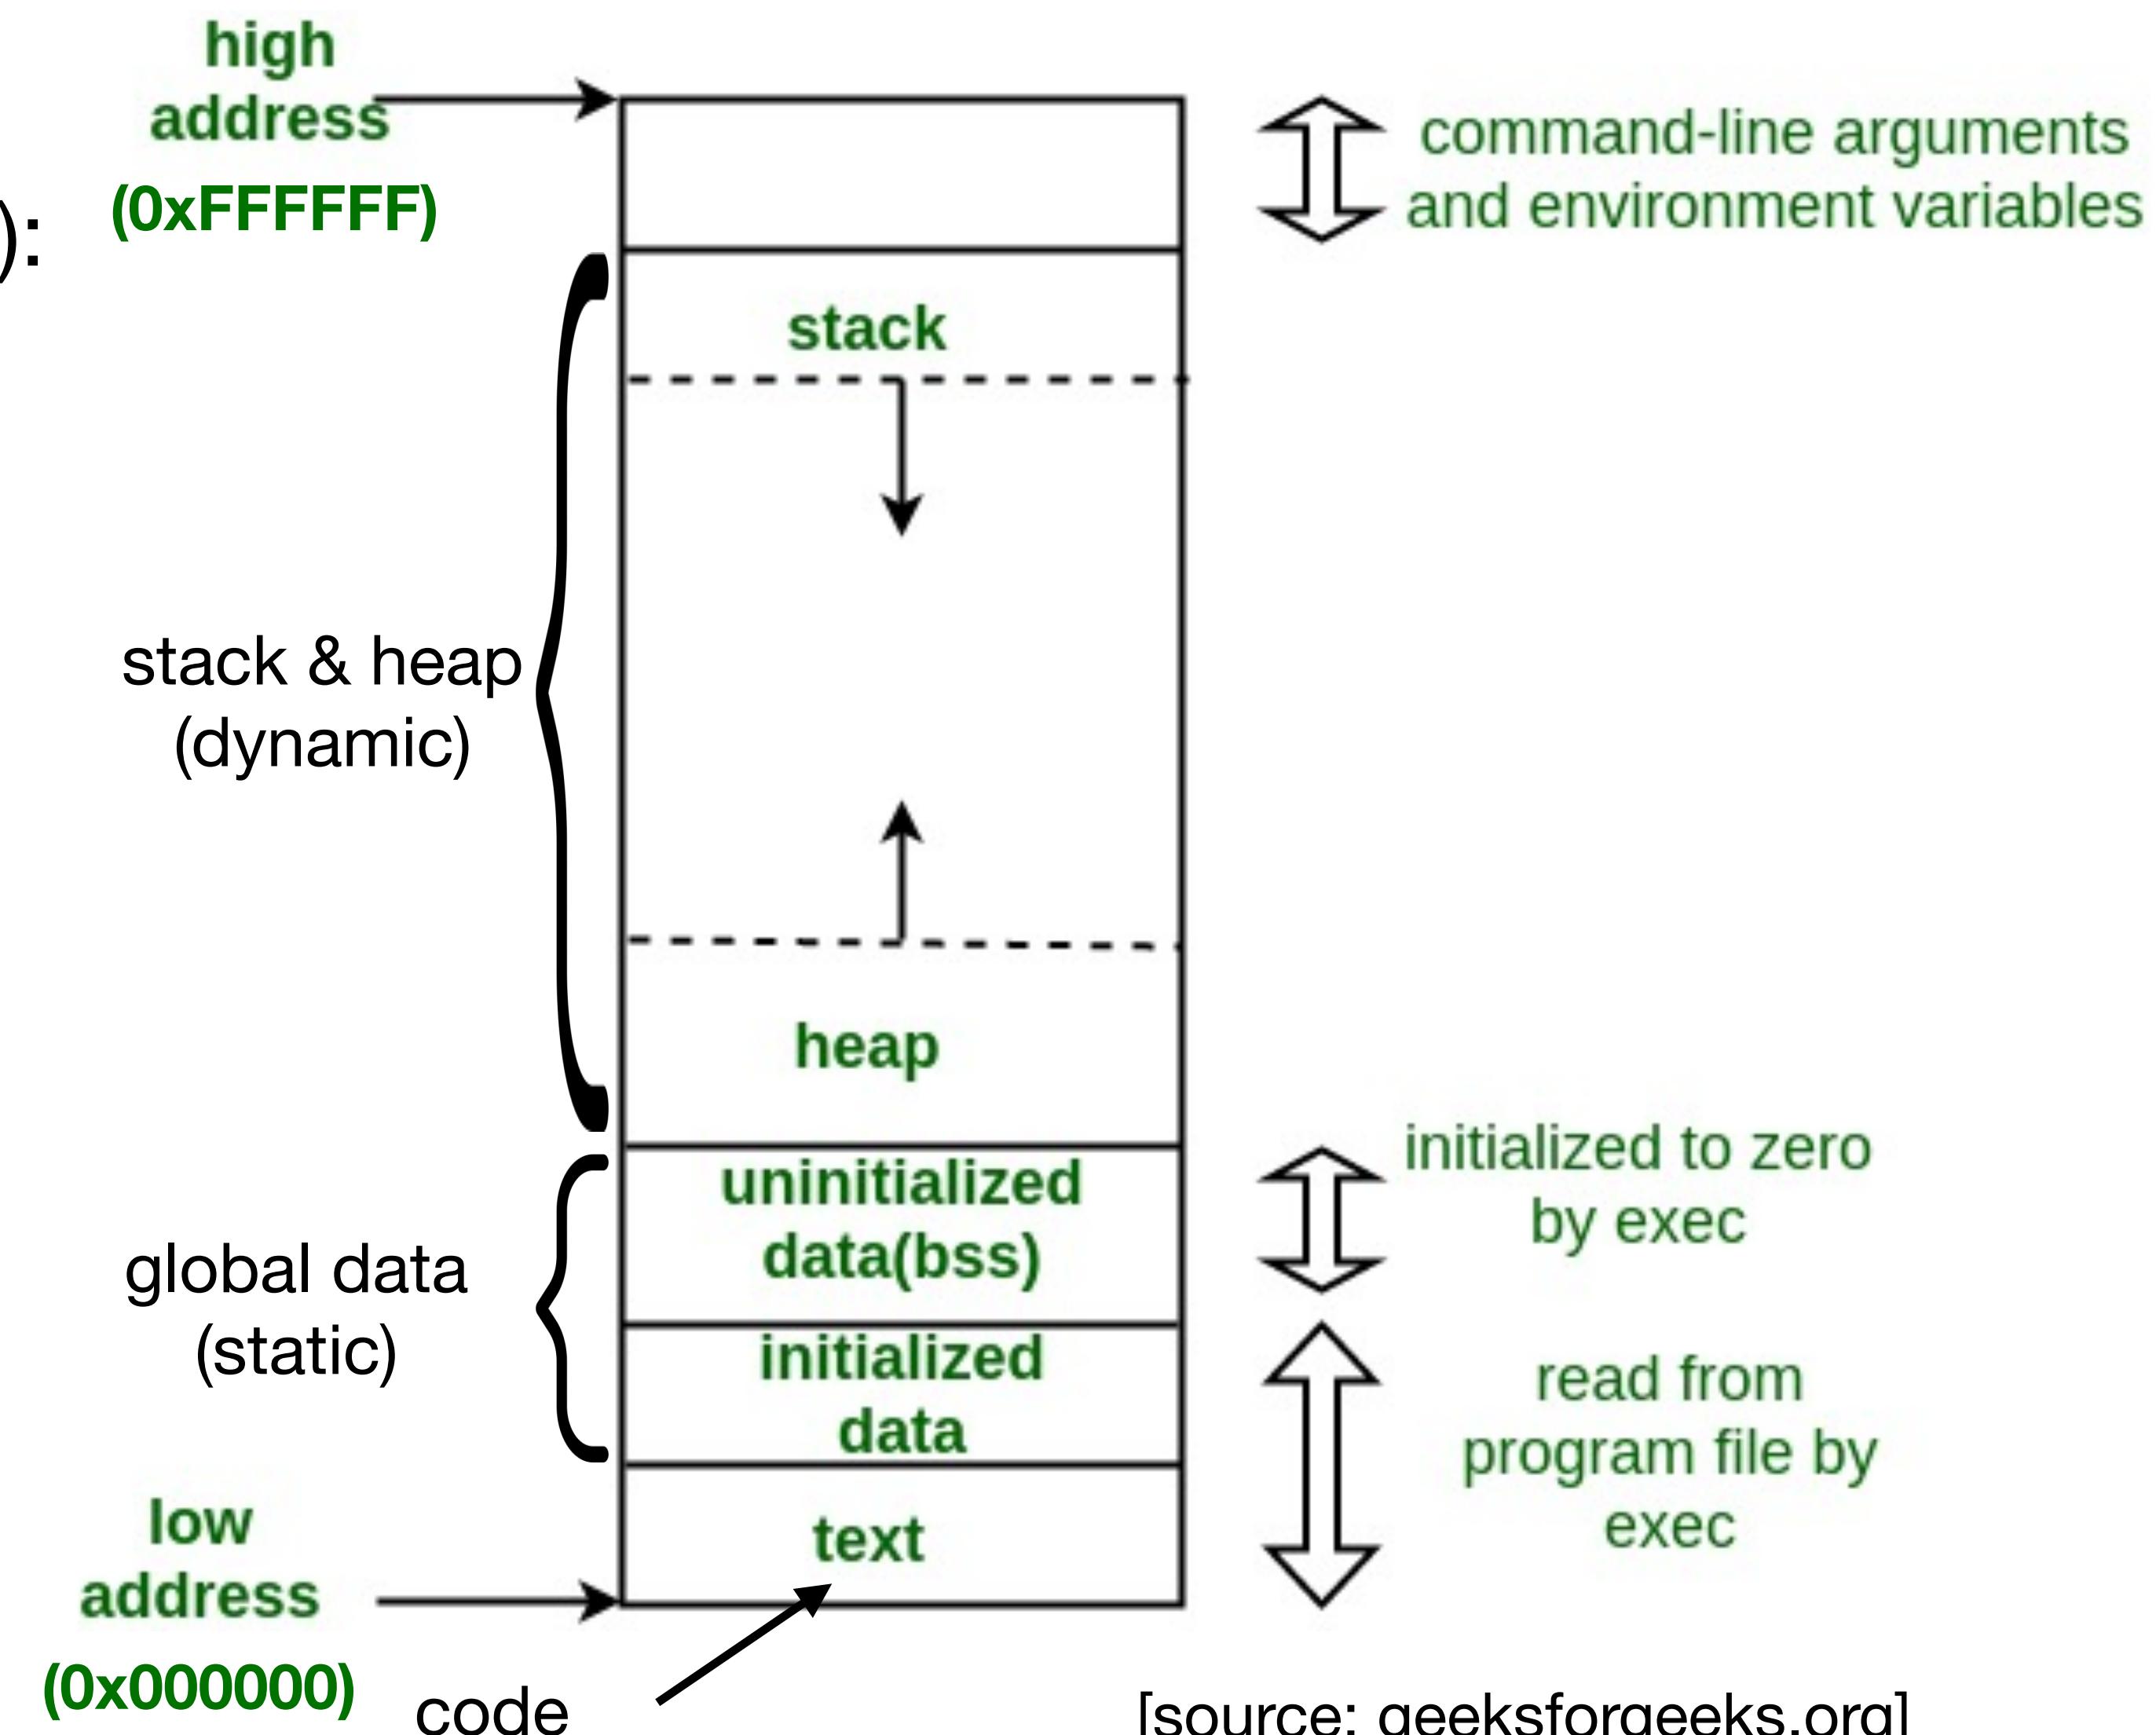
# Memory Model

- When we program in a low-level language (e.g., C)
- There is a memory management convention
- That convention is established by the HW architecture and the OS
  - We will assume Intel x86



# Memory Model

- Memory allocated automatically (**stack**):
  - return address
  - global variables
  - function/procedure arguments
- Memory allocated “on demand” (**heap**)
- Memory w/ global/static data (**data**)
- Memory w/ code (**text**)



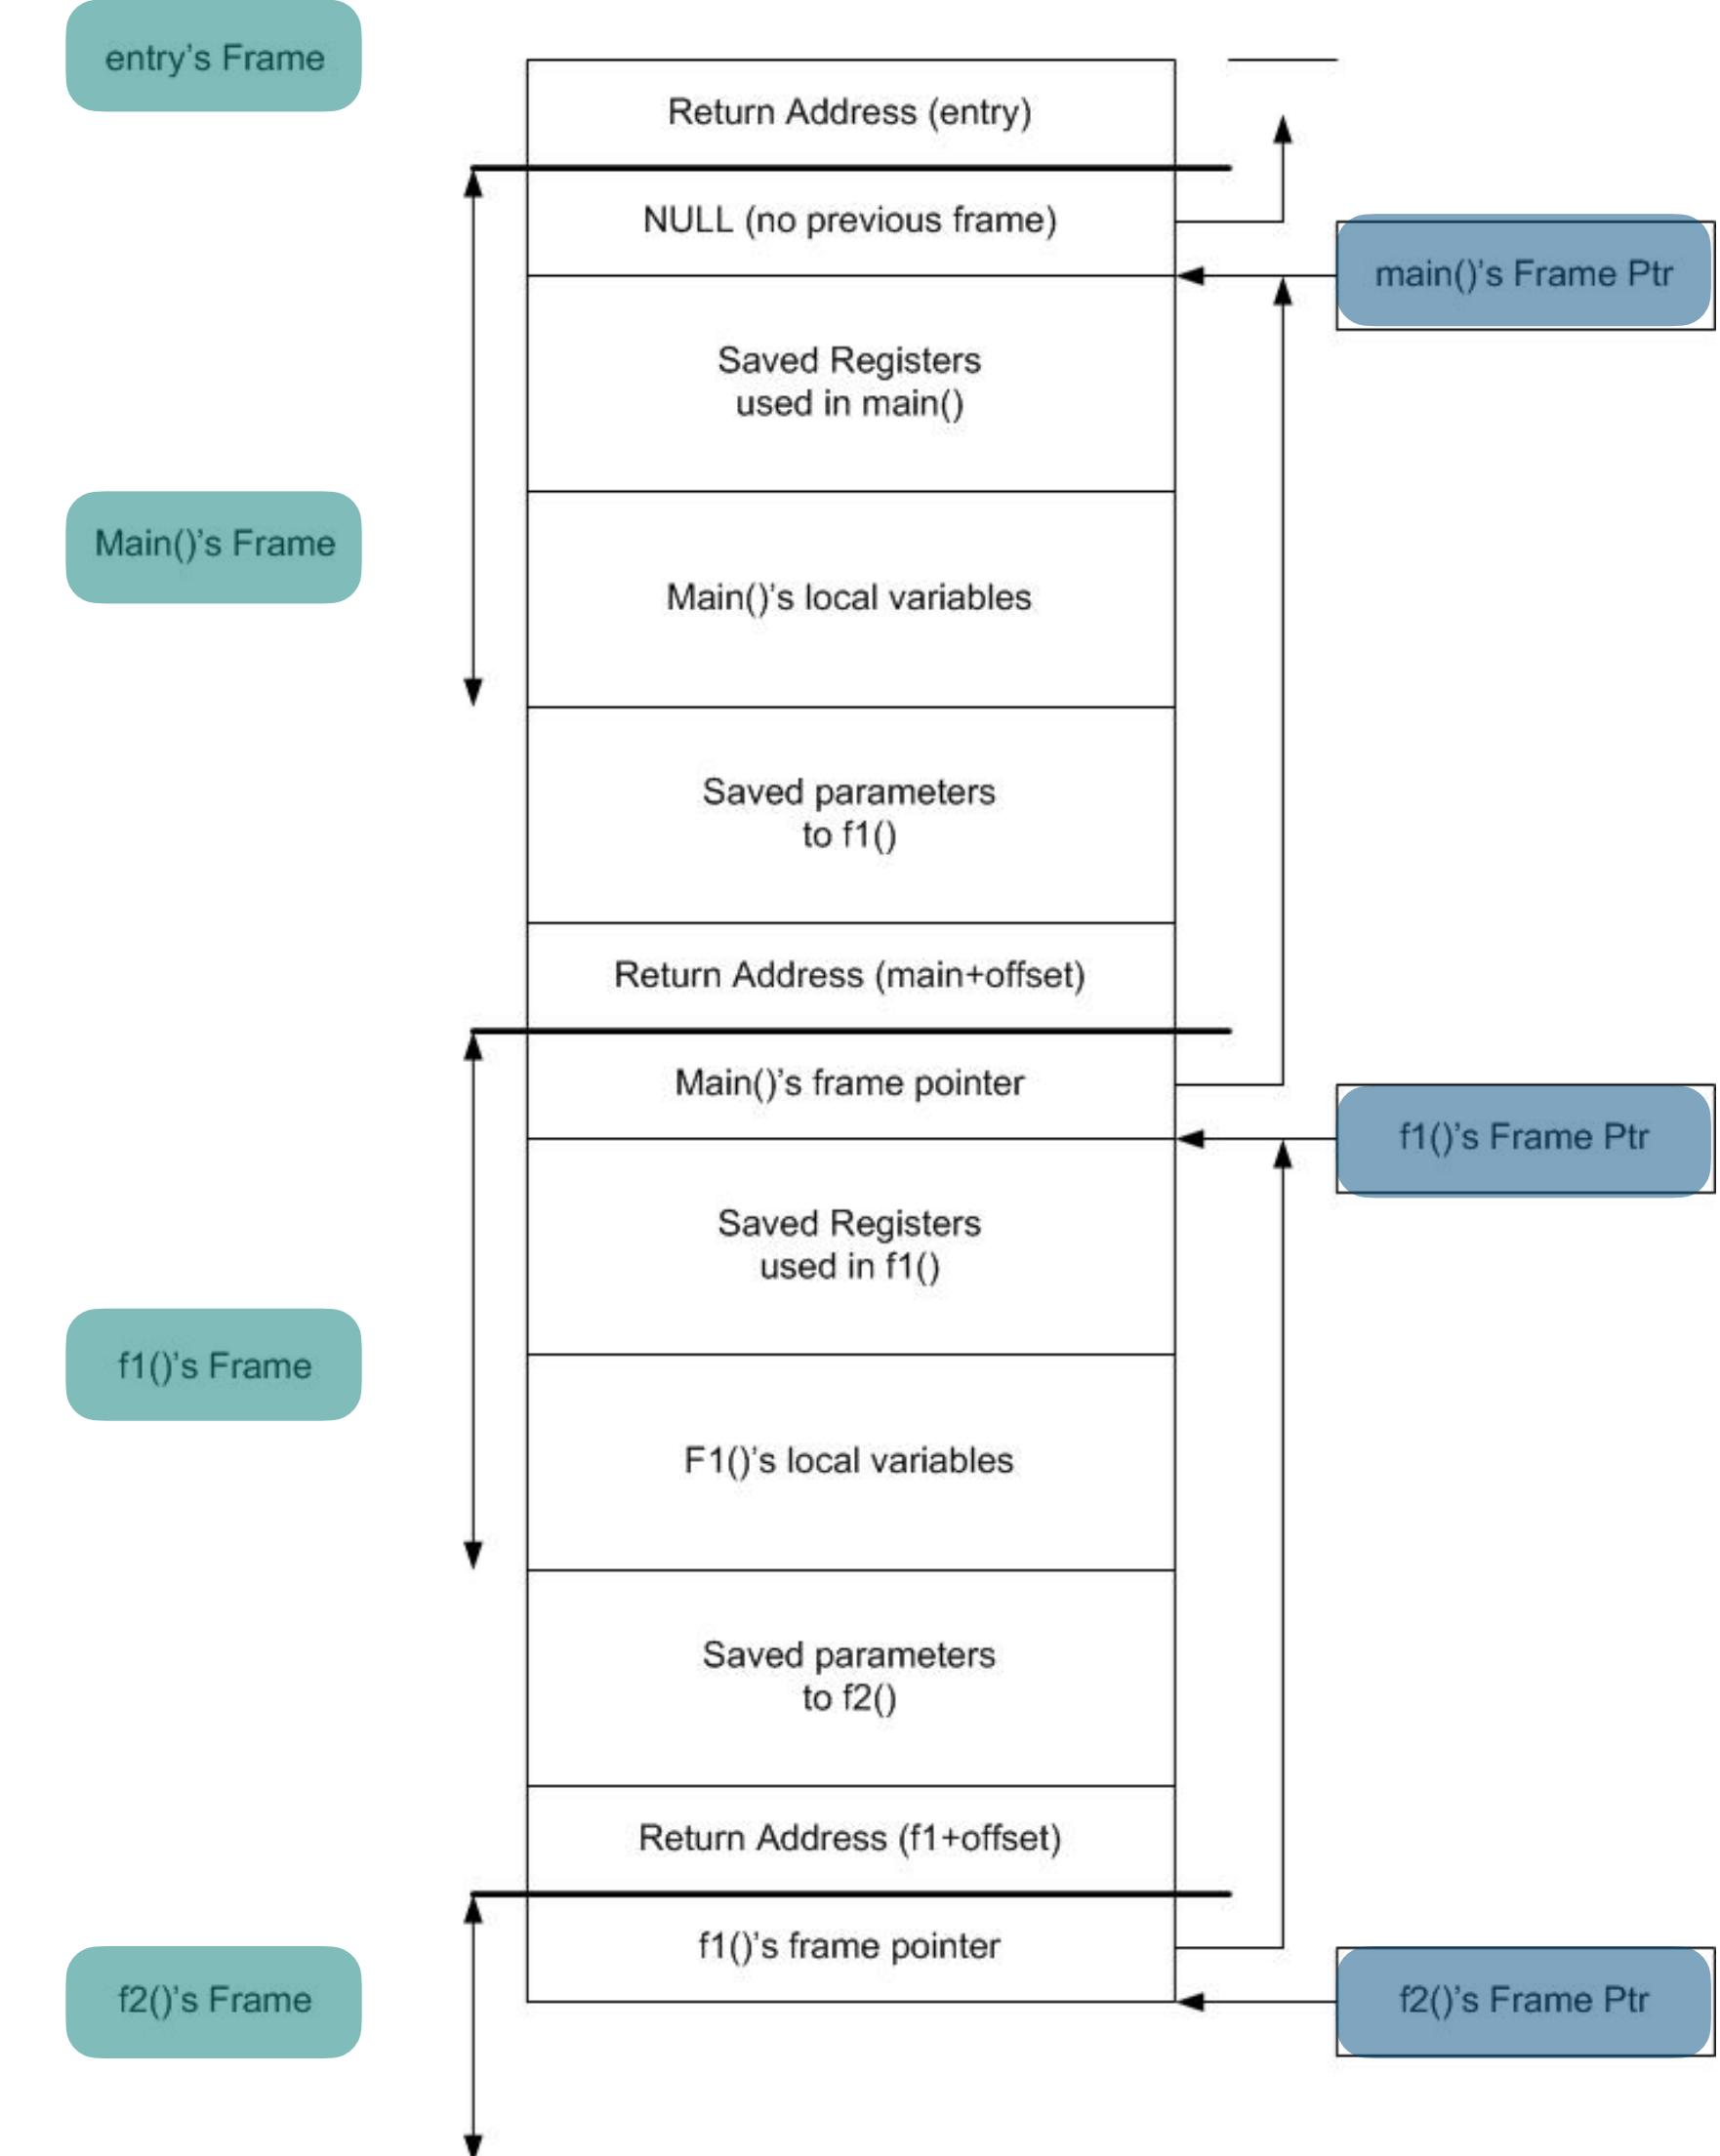
# Stealing control

- Occurs when an **external input** leads the **program** to break the convention:
  - Replacing the expected sequence of instructions for a sequence with instruction controlled by the attacker
    1. Writing the malicious sequence of instructions to memory
    2. Executing the malicious sequence of instructions from memory
  - May occur in the stack, in the heap, during system calls, etc
  - Much harder than 30 years ago
  - But still very common

# **Smashing the Stack (buffer overflows)**

# Understanding the Stack

- Each function works under a local context called a **stack frame**:
  - Remember: stack grows downwards
  - **Frame pointer** stored in a registry (e.g., ebp) points to the start of that region (local variables)
  - **Stack pointer** stored in a registry (e.g., esp) points to the end of that region
  - Contains: parameters, data stored for the return, local variables, other temporary data.
- Main(...) → f1(...) → f2(...)

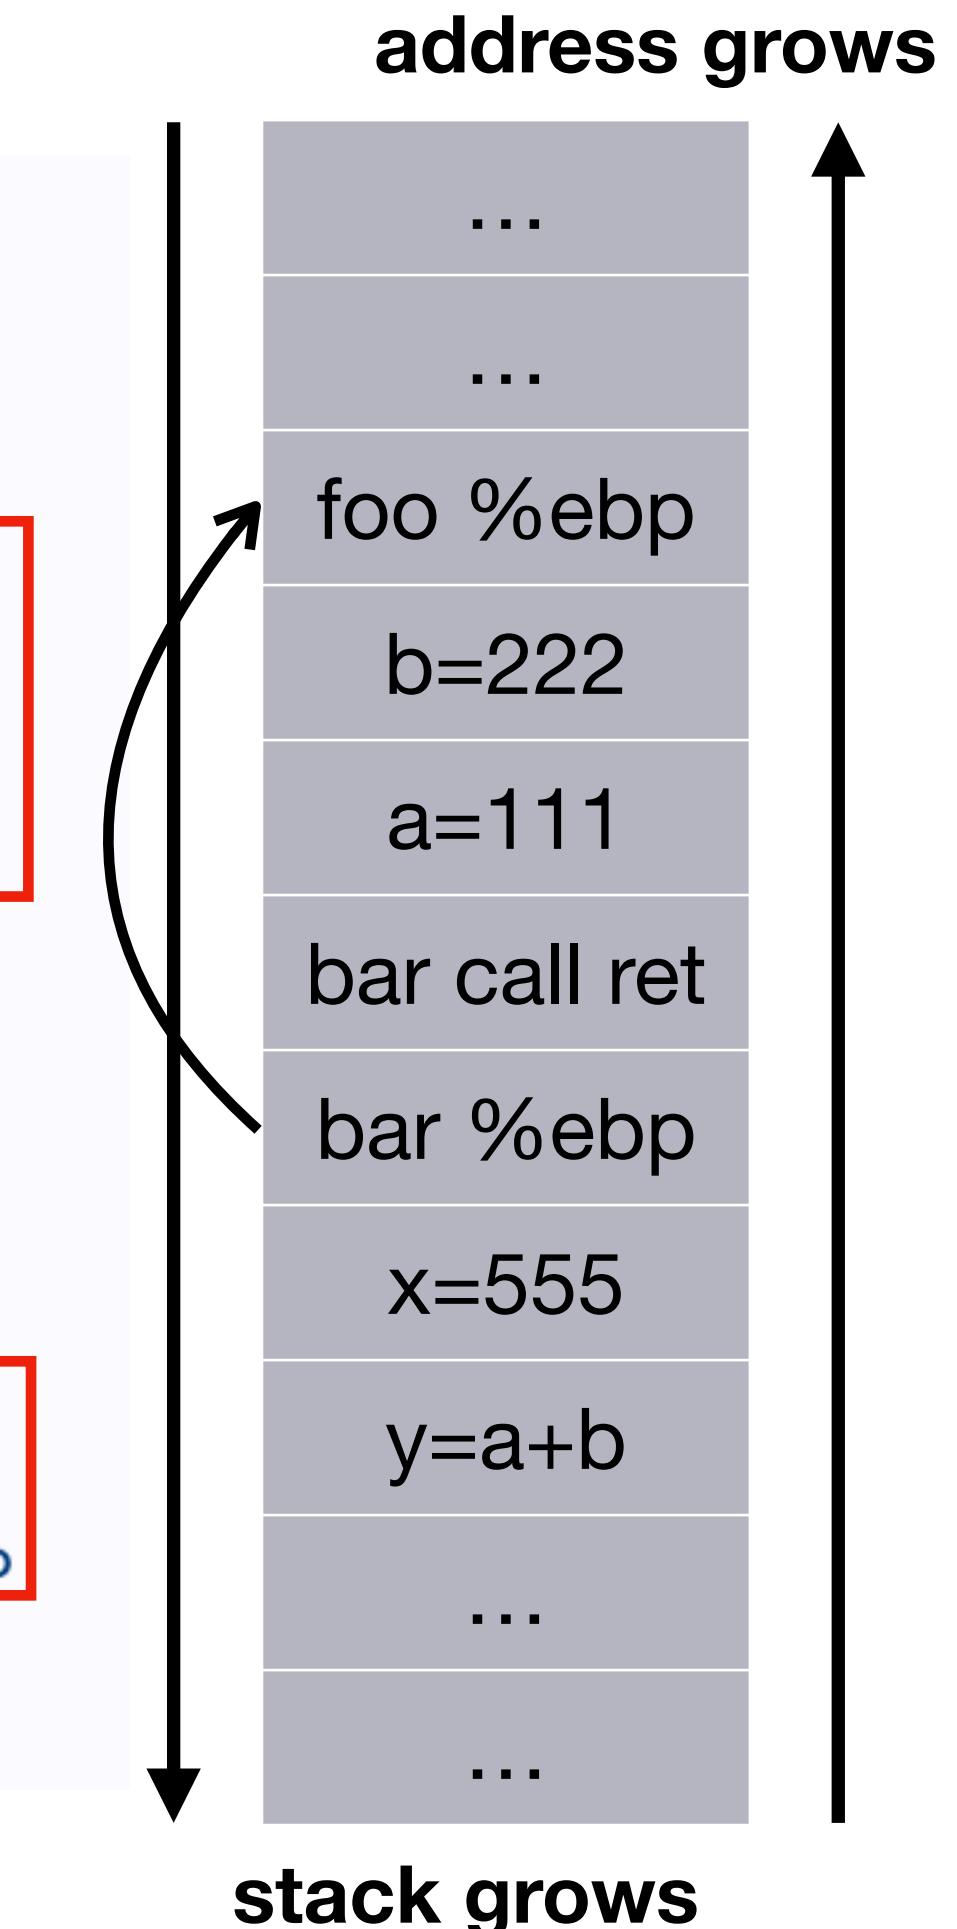


# Understanding the Stack

```
void bar(int a, int b)
{
    int x, y;
    x = 555;
    y = a+b;
}

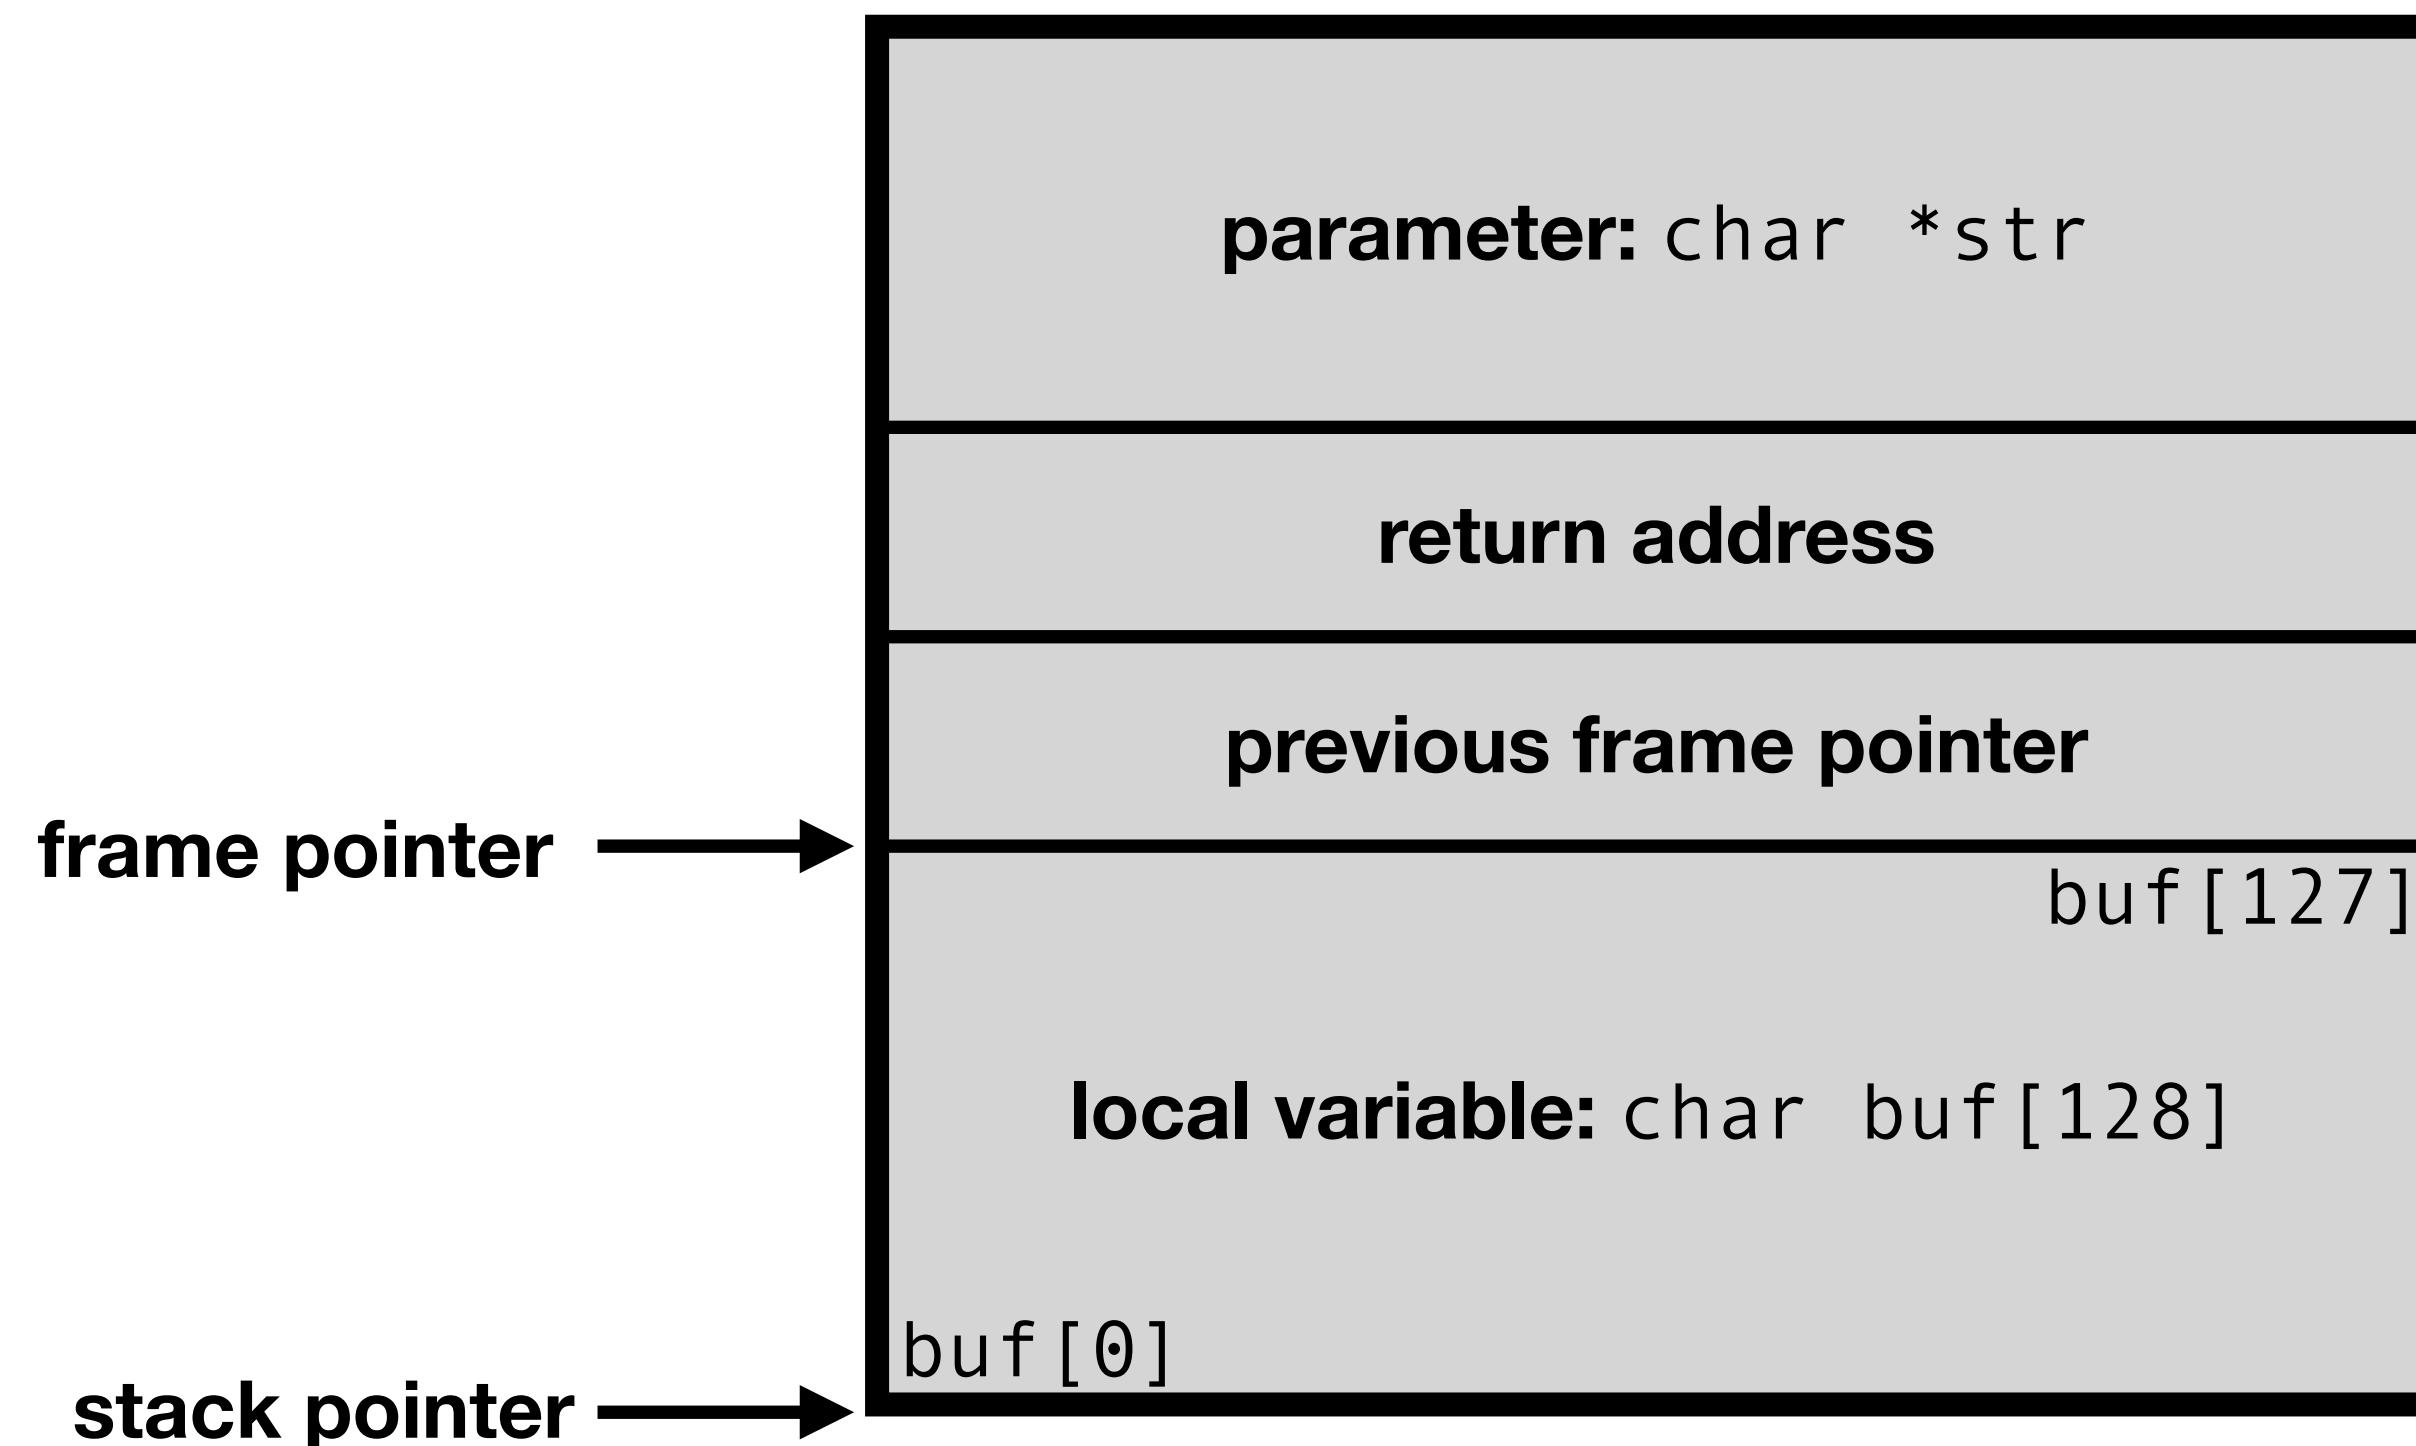
void foo(void) {
    bar(111,222);
}
```

bar: # ----- start of the function bar()  
pushl %ebp # save the incoming frame pointer  
movl %esp, %ebp # set the frame pointer to the current top of stack  
subl \$16, %esp # increase the stack by 16 bytes (stacks grow down)  
movl \$555, -4(%ebp) # x=555 a is located at [ebp-4]  
movl 12(%ebp), %eax # 12(%ebp) is [ebp+12], which is the second parameter  
movl 8(%ebp), %edx # 8(%ebp) is [ebp+8], which is the first parameter  
addl %edx, %eax # add them  
movl %eax, -8(%ebp) # store the result in y  
leave #  
ret #  
:  
# ----- start of the function foo()  
pushl %ebp # save the current frame pointer  
movl %esp, %ebp # set the frame pointer to the current top of the stack  
subl \$8, %esp # increase the stack by 8 bytes (stacks grow down)  
movl \$222, 4(%esp) # this is effectively pushing 222 on the stack  
movl \$111, (%esp) # this is effectively pushing 111 on the stack  
call bar # call = push the instruction pointer on the stack and branch to foo  
leave # done  
ret #



# Stack Smashing

- When entering func, the stack frame looks like:

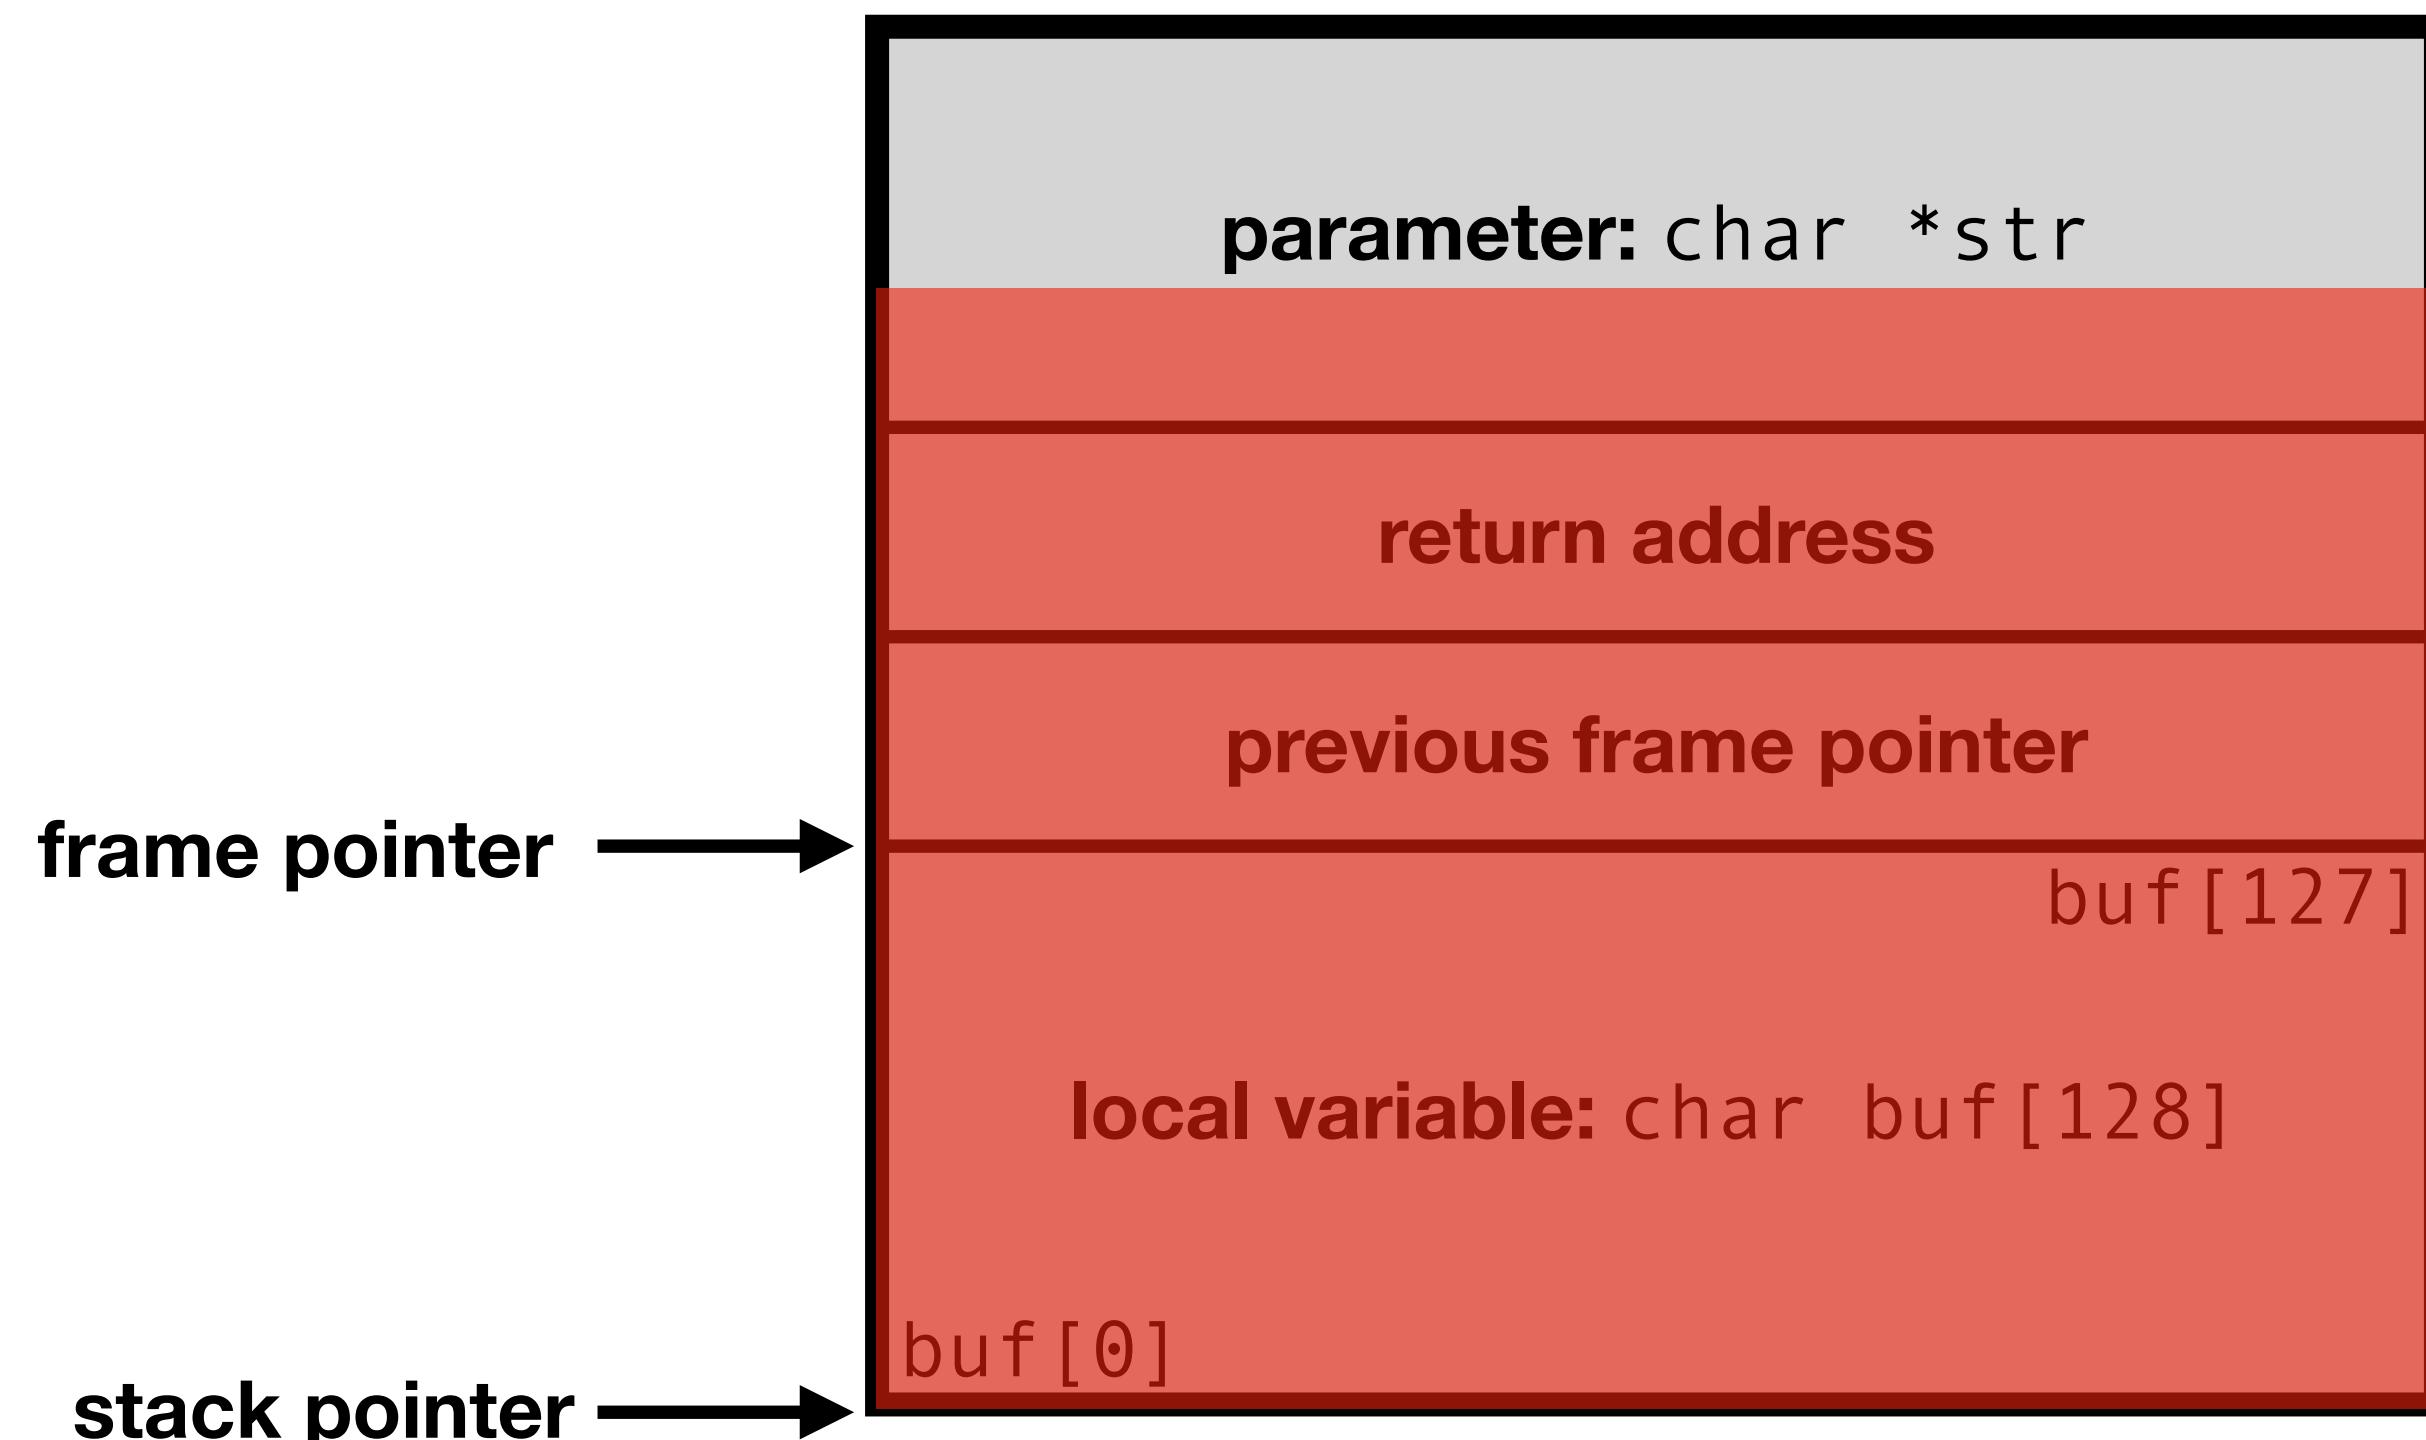


```
void func(char *str) {  
    char buf[128];  
    strcpy(buf, str);  
    user(buf);  
}
```

- What does strcpy do?
- What if str has size > 128?
- Crash? Undefined?

# Stack Smashing

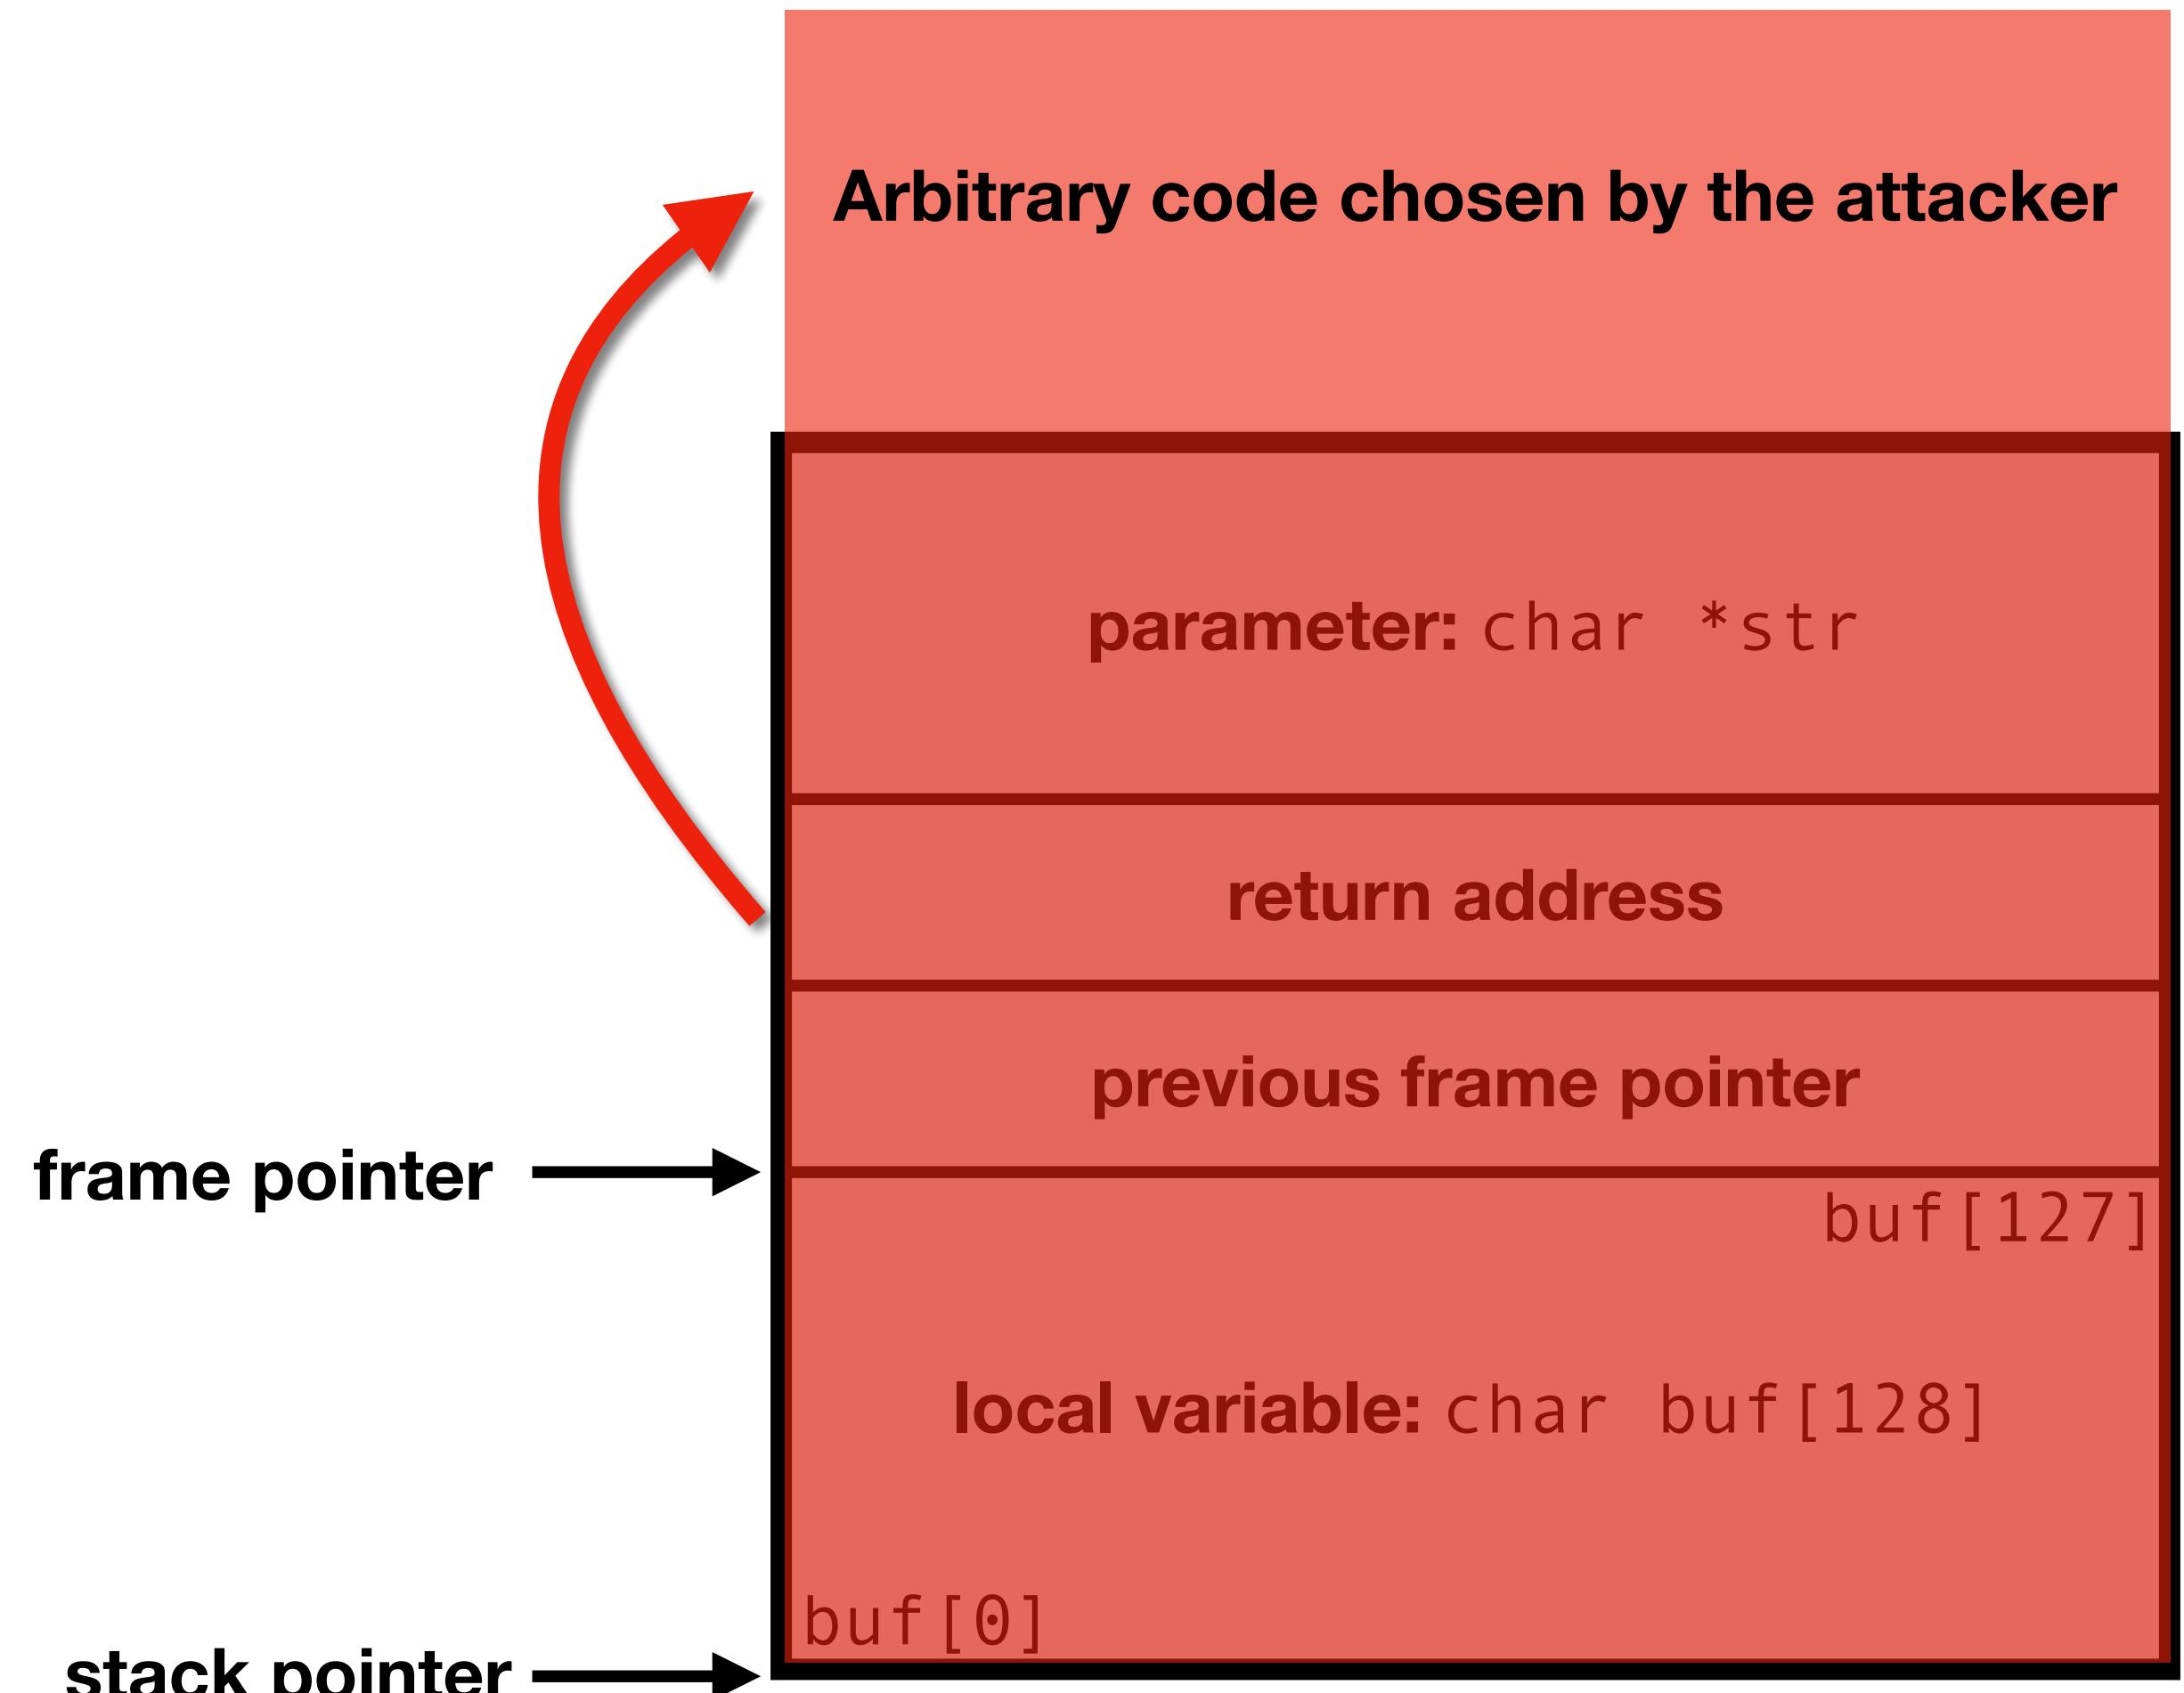
- The function writes to increasing memory positions:



```
void func(char *str) {  
    char buf[128];  
    strcpy(buf, str);  
    user(buf);  
}
```

- What happens when the return address is accessed when the function returns?

# Stack Smashing



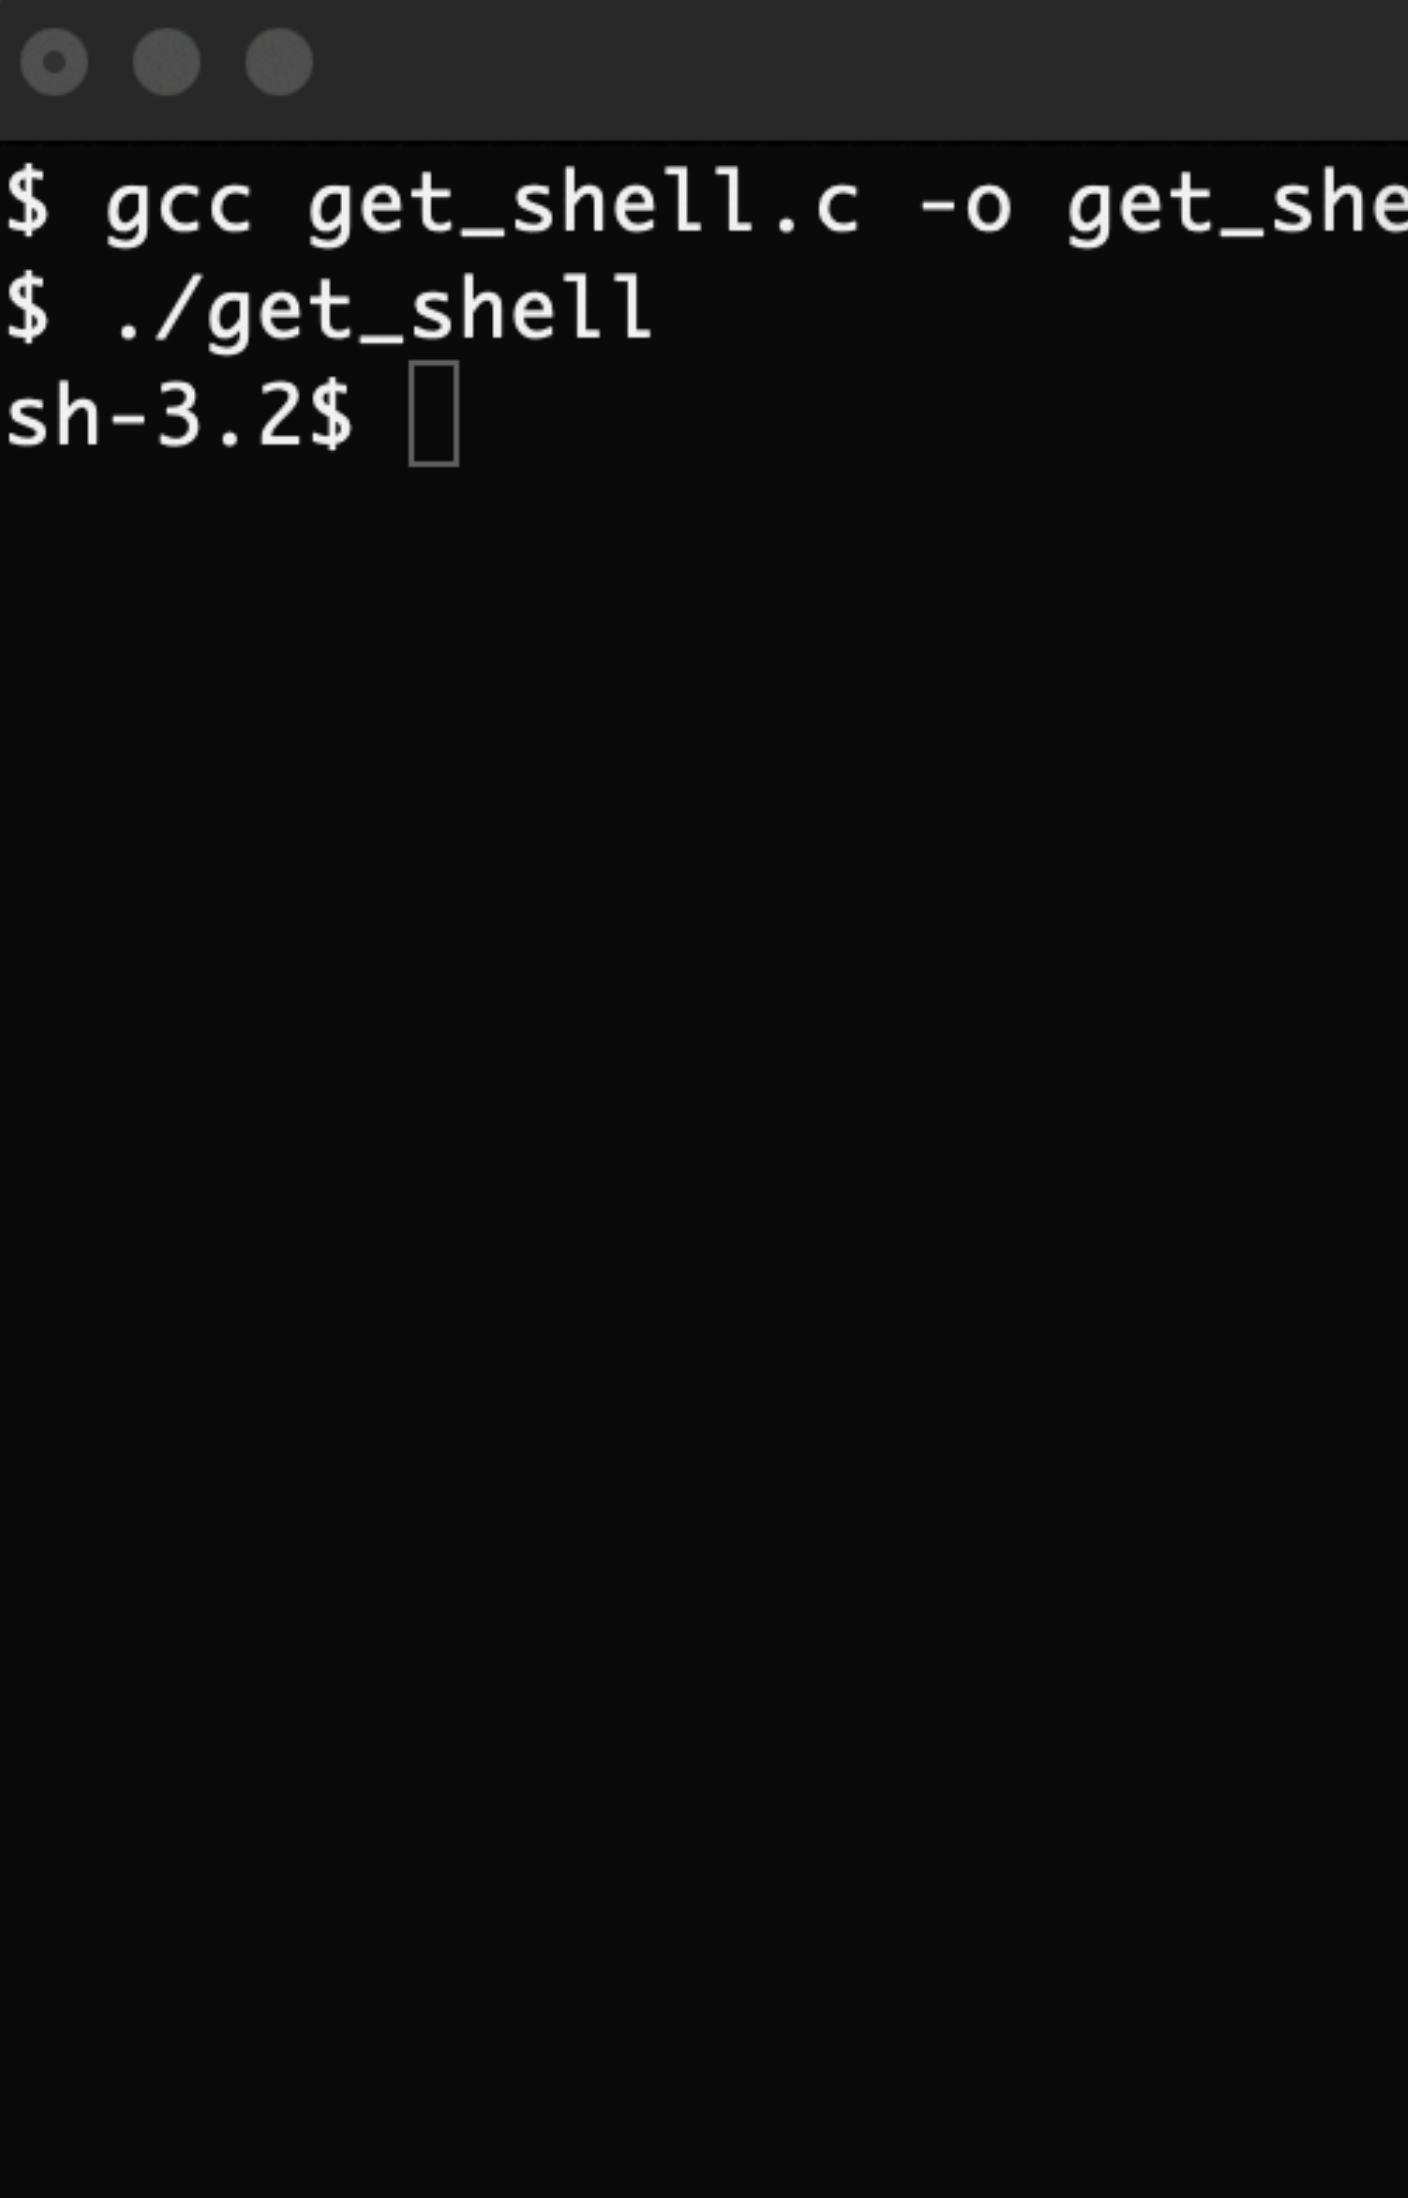
```
void func(char *str) {  
    char buf[128];  
    strcpy(buf, str);  
    user(buf);  
}
```

- If **str** comes from **outside the security perimeter**:
  - attacker may fill stack with code
  - replace the return address
  - execute arbitrary code



**There is no separation between code and data!**

# Which code to inject?



The terminal window shows the following commands:

```
$ gcc get_shell.c -o get_shell
$ ./get_shell
sh-3.2$
```

A code editor window titled "get\_shell.c" displays the C source code:

```
1 #include <unistd.h>
2
3 void get_shell() {
4     char *argv[2];
5     char *envp[1];
6     argv[0] = "/bin/sh";
7     argv[1] = NULL;
8     envp[0] = NULL;
9     execve(argv[0], argv, envp);
10 }
11
12 int main() {
13     get_shell();
14 }
```

# The Holy Grail of RCE

- “Shell code” is a sequence of instructions that executes shell commands:



- written in Assembly code
- preceded by some NOP operations (because its memory address may vary)
- injected on memory to be executed on return ⇒ buffer overflow

```
/*
Title: Linux x86 - execve("/bin/bash", ["/bin/bash", "-p"], NULL) - 33 bytes
Author: Jonathan Salwan
Mail: submit@shell-storm.org
Web: http://www.shell-storm.org

!Database of Shellcodes http://www.shell-storm.org/shellcode/

sh sets (euid, egid) to (uid, gid) if -p not supplied and uid < 100
Read more: http://www.faqs.org/faqs/unix-faq/shell/bash/#ixzz0mzPmJC49

assembly of section .text:

08048054 <.text>:
08048054: 6a 0b          push  $0xb
08048056: 58              pop   %eax
08048057: 99              cltd
08048058: 52              push  %edx
08048059: 66 68 2d 70    pushw $0x702d
0804805d: 89 e1          mov   %esp,%ecx
0804805f: 52              push  %edx
08048060: 6a 68          push  $0x68
08048062: 68 2f 62 61 73 push  $0x7361622f
08048067: 68 2f 62 69 6e push  $0x6e69622f
0804806c: 89 e3          mov   %esp,%ebx
0804806e: 52              push  %edx
0804806f: 51              push  %ecx
08048070: 53              push  %ebx
08048071: 89 e1          mov   %esp,%ecx
08048073: cd 80          int   $0x80

*/
#include <stdio.h>

char shellcode[] = "\x6a\x0b\x58\x99\x52\x66\x68\x2d\x70"
"\x89\xe1\x52\x6a\x68\x68\x2f\x62\x61"
"\x73\x68\x2f\x62\x69\x6e\x89\xe3\x52"
"\x51\x53\x89\xe1\xcd\x80";

int main(int argc, char *argv[])
{
    printf(stdout,"Length: %d\n",strlen(shellcode));
    (*(void(*)()) shellcode)();
}
```



# The devil is in the details

- Building a working **exploit** requires a lot of careful “fine tuning”
  - E.g., the codification of the program as an array of bytes cannot contain '\0' (why?)
- A robust exploit will work consistently in arbitrary executions of the vulnerable code in similar platforms
- Many variants/tricks are used to achieve such robustness
- E.g., to avoid having to guess the specific address where the *shellcode* is stored:
  - Place the *shellcode* in the stack immediately above the return address
  - Point the return address to a memory address with the instruction `jmp sp`

# Brief history [wikipedia]

- The concept is documented since 1972!
- First hostile use documented in 1988 (Morris worm)
- Generalised knowledge:
  - "Smashing the stack for fun and profit", Elias Levy/Aleph One, 1996
- Many famous/infamous examples, including both funny and disruptive ones:
  - Wii, PS2, XBox jailbreaks



[source: GameBoy Pokemon Red]  
<https://youtu.be/Nh3X8zSmTwY>

# fingerd: Morris Worm (1988)

Finger Daemon Buffer Overflow

Vulnerability Description

Brief Description: The *finger(1)* daemon is vulnerable to a buffer overrun attack, which allows a network entity to connect to the *fingerd(8)* port and **get a root shell**.

Detailed Description: *Fingerd* is a daemon that responds to requests for a listing of current users, or specific information about a particular user. It reads its input from the network, and sends its output to the network. On many systems, it ran as the *superuser* or some other privileged user. The daemon, *fingerd* uses *gets(3)* to read the data from the client. As *gets* does no bounds checking on its argument, which is an array of 512 bytes and is allocated on the stack, a longer input message will overwrite the end of the stack, changing the return address. If the appropriate code is loaded into the buffer, that code can be executed with the privileges of the *fingerd* daemon.

Component(s): finger, fingerd

Version(s): Versions before Nov. 6, 1989.

Operating System(s): All flavors of the UNIX operating system.

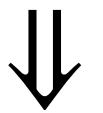
[https://minnie.tuhs.org/cgi-bin/utree.pl?  
file=4.3BSD/usr/src/etc/fingerd.c](https://minnie.tuhs.org/cgi-bin/utree.pl?file=4.3BSD/usr/src/etc/fingerd.c)



Student @ Cornell



Felon @ US CFFA



Professor @ MIT



```
main(argc, argv)
    char *argv[];
{
    register char *sp;
    char line[512];
    struct sockaddr_in sin;
    int i, p[2], pid, status;
    FILE *fp;
    char *av[4];

    i = sizeof (sin);
    if (getpeername(0, &sin, &i) < 0)
        fatal(argv[0], "getpeername");
    line[0] = '\0';
    gets(line);
```

“payload-free”  
worm that just  
disseminated itself

# Why so common?

- String/buffer manipulation via libraries such as libc
- Many functions are outright unsafe:
  - Do not guarantee reads/writes in designated memory regions
  - `strcpy`, `strcat`, `gets`, `scanf`
- Other functions may check bounds (e.g. `strncpy`) but do not guarantee that string is properly terminated
- Or simply code implemented from scratch with similar problems:
  - It is often assumed that **input** comes from a **trusted source** ⇒ 
  - **Memory management vulnerabilities are really subtle!**



# libpng (2004)

## libpng png\_handle\_tRNS() Buffer Overflow May Let Remote Users Execute Arbitrary Code

SecurityTracker Alert ID: 1011848

SecurityTracker URL: <http://securitytracker.com/id/1011848>

CVE Reference: [GENERIC-MAP-NOMATCH](#) (*Links to External Site*)

Updated: Oct 21 2004

Original Entry Date: Oct 21 2004

Impact: [Execution of arbitrary code via network](#), [User access via network](#)

Fix Available: Yes Vendor Confirmed: Yes

Description: A buffer overflow vulnerability was reported in libpng. A remote user may be able to execute arbitrary code on the target system.

Debian reported that an image with certain height or width values can trigger an overflow in the png\_handle\_tRNS() function in pngutil.c.

Impact: A remote user may be able to execute arbitrary code on a system that uses libpng.

Solution: Versions 1.0.17 and 1.2.7 include the fix and are available at:

[http://sourceforge.net/project/showfiles.php?group\\_id=5624](http://sourceforge.net/project/showfiles.php?group_id=5624)

[Editor's note: The fix may have been introduced into earlier versions.]

Vendor URL: [libpng.sourceforge.net](#) (*Links to External Site*)

Cause: Boundary error

Underlying OS: [Linux \(Any\)](#), [UNIX \(Any\)](#)



### Lack of input validation

```
if (!(png_ptr->mode & PNG_HAVE_PLTE))
{
    /* Should be an error, but we can cope with it */
    png_warning(png_ptr, "Missing PLTE before tRNS");
}
else if (length > (png_uint_32)png_ptr->num_palette)
{
    png_warning(png_ptr, "Incorrect tRNS chunk length");
    png_crc_finish(png_ptr, length);
    return;
}
```

We can see, if the first warning condition is hit, the length check is missed due to the use of an "else if".

# libpng (2004)

## CVE Details

The ultimate security vulnerability datasource

[Log In](#) [Register](#)

**What's the CVSS score of your company?**

Search

(e.g.: CVE-2009-1234 or 2010-1234 or 20101234)

View CVE



[Switch to https://](#)

[Home](#)

**Browse :**

[Vendors](#)

[Products](#)

[Vulnerabilities By Date](#)

[Vulnerabilities By Type](#)

**Reports :**

[CVSS Score Report](#)

[CVSS Score Distribution](#)

**Search :**

[Vendor Search](#)

[Product Search](#)

[Version Search](#)

[Vulnerability Search](#)

[By Microsoft References](#)

**Top 50 :**

[Vendors](#)

[Vendor Cvss Scores](#)

[Products](#)

[Product Cvss Scores](#)

[Versions](#)

**Other :**

[Microsoft Bulletins](#)

[Bugtraq Entries](#)

...  
...

### Vulnerability Details : [CVE-2004-0597](#)

Multiple buffer overflows in libpng 1.2.5 and earlier, as used in multiple products, allow remote attackers to execute arbitrary code via malformed PNG images in which (1) the png\_handle\_tRNS function does not properly validate the length of transparency chunk (tRNS) data, or the (2) png\_handle\_sBIT or (3) png\_handle\_hIST functions do not perform sufficient bounds checking.

Publish Date : 2004-11-23 Last Update Date : 2018-10-12

[Collapse All](#) [Expand All](#) [Select](#) [Select&Copy](#)

▼ [Scroll To](#) ▼ [Comments](#) ▼ [External Links](#)

[Search Twitter](#) [Search YouTube](#) [Search Google](#)

#### - CVSS Scores & Vulnerability Types

CVSS Score

**10.0**

Confidentiality Impact

**Complete** (There is total information disclosure, resulting in all system files being revealed.)

Integrity Impact

**Complete** (There is a total compromise of system integrity. There is a complete loss of system protection, resulting in the entire system being compromised.)

Availability Impact

**Complete** (There is a total shutdown of the affected resource. The attacker can render the resource completely unavailable.)

Access Complexity

**Low** (Specialized access conditions or extenuating circumstances do not exist. Very little knowledge or skill is required to exploit. )

Authentication

**Not required** (Authentication is not required to exploit the vulnerability.)

Gained Access

**None**

Vulnerability Type(s)

Execute Code Overflow

CWE ID

CWE id is not defined for this vulnerability

...  
...

# Off-by-1 is enough (realpath, 2003)

```
/*
 * Join the two strings together, ensuring that the right thing
 * happens if the last component is empty, or the dirname is root.
 */
if (resolved[0] == '/' && resolved[1] == '\0')
    rootd = 1;
else
    rootd = 0;

if (*wbuf) {
    if (strlen(resolved) + strlen(wbuf) + rootd + 1 > MAXPATHLEN) {
        errno = ENAMETOOLONG;
        goto err1;
    }
    if (rootd == 0)
        (void)strcat(resolved, "/");
    (void)strcat(resolved, wbuf);
}
```

source: <https://www.securityfocus.com/bid/8315/info>

## Details:

=====

An off-by-one bug exists in `fb_realpath()` function. An overflow occurs when the length of a constructed path is equal to the `MAXPATHLEN+1` characters while the size of the buffer is `MAXPATHLEN` characters only. The overflowed buffer lies on the stack.

The bug results from misuse of `rootd` variable in the calculation of length of a concatenated string:

<https://www.exploit-db.com/exploits/22974> 

The '`realpath()`' function is a C-library procedure to resolve the canonical, absolute pathname of a file based on a path that may contain values such as '/', './', '../', or symbolic links. A vulnerability that was reported to affect the implementation of '`realpath()`' in WU-FTPD has lead to the discovery that at least one implementation of the C library is also vulnerable. FreeBSD has announced that the off-by-one stack- buffer-overflow vulnerability is present in their libc. Other systems are also likely vulnerable.

Reportedly, this vulnerability has been successfully exploited against WU-FTPD to execute arbitrary instructions.

NOTE: Patching the C library alone may not remove all instances of this vulnerability. Statically linked programs may need to be rebuilt with a patched version of the C library. Also, some applications may implement their own version of '`realpath()`'. These applications would require their own patches. FreeBSD has published a large list of applications that use '`realpath()`'. Administrators of FreeBSD and other systems are urged to review it. For more information, see the advisory 'FreeBSD-SA-03:08.realpath'.

# The potential was there! (1998)

## The poisoned NUL byte

*From:* okir () MONAD SWB DE (Olaf Kirch)

*Date:* Wed, 14 Oct 1998 11:42:46 +0200

---

Summary: you can exploit a single-byte buffer overrun to gain root privs.

The interesting thing about this overrun is that it was by just a single byte. And yes, it not just crashed the process, it provided a root shell. It took me a while to figure out, but what it boils down to is this:

At the beginning of the function, realpath copies the argument (1024 bytes) To a local buffer (sized MAXPATHLEN, i.e. 1024 bytes). Thus, the terminating 0 byte of the string gets scribbled over the next byte, which happens to be the **lowest byte of %ebp**, the frame pointer of the calling function. At function entry, its value was 0xbffff3ec. After the strcpy, it becomes 0xbffff300.

During the remainder of realpath(), nothing exciting happens, but when the function returns, %ebp is restored from stack, which effectively shifts down the calling function's stack frame by 0xec bytes.

The calling function now does a few things with local data, dereferences some pointers (by sheer dumb luck these pointers contain random but valid addresses), and returns, restoring the %esp and %ebp registers from stack. With the stack having shifted down 0xec bytes, it picks up the return address from the local buffer containing the exploit code...

# And still exists (2014)

News and updates from the Project Zero team at Google

Monday, August 25, 2014

## The poisoned NUL byte, 2014 edition

Posted by Chris Evans, Exploit Writer Underling to Tavis Ormandy

Back in [this 1998 post to the Bugtraq mailing list](#), Olaf Kirch outlined an attack he called "The poisoned NUL byte". It was an off-by-one error leading to writing a NUL byte outside the bounds of the current stack frame. On i386 systems, this would clobber the least significant byte (LSB) of the "saved %ebp", leading eventually to code execution. Back at the time, people were surprised and horrified that such a minor error and corruption could lead to the compromise of a process.

Fast forward to 2014. Well over a month ago, Tavis Ormandy of Project Zero [disclosed a glibc NUL byte off-by-one overwrite into the heap](#). Initial reaction was [skepticism about the exploitability of the bug](#), on account of the malloc metadata hardening in glibc. In situations like this, the Project Zero culture is to sometimes "wargame" the situation. geohot quickly coded up a challenge and we were able to gain code execution. Details are captured [in our public bug](#). This bug contains analysis of a few different possibilities arising from an off-by-one NUL overwrite, a solution to the wargame (with comments), and of course a couple of different variants of a full exploit (with comments) for a local Linux privilege escalation.

Inspired by the success of the wargame, I decided to try and exploit a real piece of software. I chose the "pkexec" setuid binary as used by Tavis to demonstrate the bug. The goal is to attain root privilege escalation. Outside of the wargame environment, it turns out that there are a series of very onerous constraints that make exploitation hard. I did manage to get an exploit working, though, so read on to see how.

Project Zero < 

 [googleprojectzero.blogspot.com](https://googleprojectzero.blogspot.com)

Project Zero is a team of security analysts employed by Google tasked with finding zero-day vulnerabilities. It was announced on 15 July 2014. [Wikipedia](#)

**Founded:** July 15, 2014

# Are there no alternatives?

- Why not use higher-level languages with better memory safety guarantees? Java, JavaScript, Python, Rust, Haskell, etc
  - They often compile to and/or use C/LLVM libraries
- Virtual machines, sandboxes, browsers, OSes still have to be programmed using low-level languages
- In specialised domains (e.g., cryptographic code), humans are still better at aggressively optimising code than compilers
- Many attacks chain together various of these vulnerabilities
- As we will keep seeing, **input is always potentially malicious**

# Acknowledgements

- This lecture's slides have been inspired by the following lectures:
  - CSE127: Stack Buffer Overflows + Buffer Overflows and Stack Smashing
  - CS155: Basic Control Hijacking Attacks
  - CS343: Control Flow + Control Flow II + Control Flow III