

# Computer Security Foundations

## Week 8: Symmetric Encryption

Bernardo Portela

L.EIC - 24

## Context

People usually think of *encryption* when they hear *cryptography*

... which is not unreasonable. But it is only one of many cryptographic techniques

## Context

People usually think of *encryption* when they hear *cryptography*

... which is not unreasonable. But it is only one of many cryptographic techniques

Encryption guarantees *confidentiality*, but real-world applications often require other guarantees to be considered secure systems

- Authenticity, non-repudiation, unpredictability, anonymity, ...

## Context

People usually think of *encryption* when they hear *cryptography*

... which is not unreasonable. But it is only one of many cryptographic techniques

Encryption guarantees *confidentiality*, but real-world applications often require other guarantees to be considered secure systems

- Authenticity, non-repudiation, unpredictability, anonymity, ...

Also, there are many kinds of encryption

- Symmetric, asymmetric, authenticated, homomorphic, ...


# What is encryption?

**Q1: What do you think encryption means?**

# What is encryption?

## Q1: What do you think encryption means?


Encryption transforms *plaintexts* into *ciphertexts* using a *key*



# What is encryption?

## Q1: What do you think encryption means?

Encryption transforms *plaintexts* into *ciphertexts* using a *key*



We will use the following notation to talk about algorithms

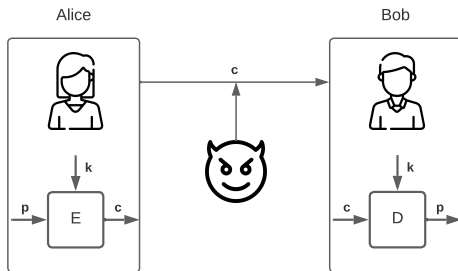
- $c \leftarrow E(k, p)$  - Encryption is (usually) randomized
  - **Q2: Why?**
- $p \leftarrow D(k, c)$  - Decryption is deterministic

We begin with **symmetric** encryption: same key on both ends

# What we talk about when we talk about Security - Part 1

## Meet Alice and Bob

- Alice wants to send a message to Bob
- The message must be secure against an attacker (the devil)






# What we talk about when we talk about Security - Part 1


## Meet Alice and Bob

- Alice wants to send a message to Bob
- The message must be secure against an attacker (the devil)



**Q: What do we mean for the encryption to be “secure”?**


## What we talk about when we talk about Security - Part 2



Suppose my message is “banana”

- Attempt #1: I don't want “*banana*” to be revealed


## What we talk about when we talk about Security - Part 2



Suppose my message is “banana”

- Attempt #1: I don't want “*banana*” to be revealed
  - But if a scheme reveals “*bananb*” I am also not happy.
- Attempt #2: I don't want any characters to be retrieved


## What we talk about when we talk about Security - Part 2



Suppose my message is “banana”

- Attempt #1: I don't want “*banana*” to be revealed
  - But if a scheme reveals “bananb” I am also not happy.
- Attempt #2: I don't want any characters to be retrieved
  - But if a scheme reveals “cbobob” I am also not happy.

## What we talk about when we talk about Security - Part 2



Suppose my message is “banana”

- Attempt #1: I don't want “*banana*” to be revealed
  - But if a scheme reveals “bananb” I am also not happy.
- Attempt #2: I don't want any characters to be retrieved
  - But if a scheme reveals “cbobob” I am also not happy.
- A more rigorous approach to define security must be taken

## Caesar Cipher


- One of the earliest known ciphers, used in Roman times

## Caesar Cipher

- One of the earliest known ciphers, used in Roman times

### Algorithm

- Take the plaintext, e.g., "banana"
- Shift the plaintext 3 letters upward
- Key is fixed, but we can choose other shift sizes as keys




# Caesar Cipher

- One of the earliest known ciphers, used in Roman times

## Algorithm

- Take the plaintext, e.g., "banana"
- Shift the plaintext 3 letters upward
- Key is fixed, but we can choose other shift sizes as keys



- **Q1: How can we decrypt?**




# Caesar Cipher

- One of the earliest known ciphers, used in Roman times

## Algorithm

- Take the plaintext, e.g., "banana"
- Shift the plaintext 3 letters upward
- Key is fixed, but we can choose other shift sizes as keys



- **Q1: How can we decrypt?**
- **Q2: Why is this cipher insecure?** *Very small key space!*

## Substitution Ciphers

- We can choose different shifts for different letters
  - E.g.  $'a' \rightarrow 'f'$ ;  $'b' \rightarrow 'a'$ ;  $'c' \rightarrow 'z'$ ; ...
- Shift is a particular class of permutations over the alphabet
  - **Q: How many permutations are there over the alphabet?**
  - **A.k.a. how large is the key space?**

## Substitution Ciphers

- We can choose different shifts for different letters
  - E.g.  $'a' \rightarrow 'f'$ ;  $'b' \rightarrow 'a'$ ;  $'c' \rightarrow 'z'$ ; ...
- Shift is a particular class of permutations over the alphabet
  - **Q: How many permutations are there over the alphabet?**
  - **A.k.a. how large is the key space?**
- $26! \approx 2^{88}$ : It's a pretty big number
- Not possible to brute force without massive investment
- Surely it will be safe... Right?

## Frequency letter attacks

**Q1: Which of these is most common in Portuguese?**

1. 'l'
2. 'a'
3. 's'
4. 'z'

## Frequency letter attacks

**Q1: Which of these is most common in Portuguese?**

1. 'l' - 2.78%
2. 'a' - 14.63%
3. 's' - 6.81%
4. 'z' - 0.47%

## Frequency letter attacks

**Q1: Which of these is most common in Portuguese?**

1. 'l' - 2.78%
2. 'a' - 14.63%
3. 's' - 6.81%
4. 'z' - 0.47%

**Q2: How can we use this to attack this encryption scheme?**

## Frequency letter attacks

**Q1: Which of these is most common in Portuguese?**

1. 'l' - 2.78%
2. 'a' - 14.63%
3. 's' - 6.81%
4. 'z' - 0.47%

**Q2: How can we use this to attack this encryption scheme?**

- Gather many ciphertexts and count the frequency of letters
- Match that frequency with the frequency of plaintext alphabet
  - With good odds, the most common letter in the ciphertexts will match the most common letter in the plaintext alphabet
- Can be done using a statistical hypothesis ( $\chi^2$ ) test

## Rotor machines

### Hebern machine (left)

- Key is the disk, encoding a substitution table
- On key press, the output is encrypted and the disk rotates

### The Enigma (right)

- Key is the initial setting of rotors by multiple rotors (3-5)
- Rotors rotate with different frequencies





## The one-time pad - Part 1

- Patent issued in 1917 by Gilbert Vernam

### Algorithm


- Lets work with 0s and 1s
- Choose a random bit string  $k \leftarrow \{0, 1\}^m$
- To encrypt, compute the bit-wise XOR of  $m$  and  $k$ :  $m \oplus k$
- To decrypt, compute the bit-wise XOR of  $c$  and  $k$ :  $c \oplus k$

## The one-time pad - Part 1


- Patent issued in 1917 by Gilbert Vernam

### Algorithm

- Lets work with 0s and 1s
- Choose a random bit string  $k \leftarrow \{0, 1\}^m$
- To encrypt, compute the bit-wise XOR of  $m$  and  $k$ :  $m \oplus k$
- To decrypt, compute the bit-wise XOR of  $c$  and  $k$ :  $c \oplus k$




## The one-time pad - Part 2



**Q1: Is this secure?**


## The one-time pad - Part 2



**Q1: Is this secure?**

- It is *perfectly secure* (as long as keys are used only once)

## The one-time pad - Part 2




**Q1: Is this secure?**

- It is *perfectly secure* (as long as keys are used only once)

**Q2: Why is this not used to encrypt everything?**

## The one-time pad - Part 2



### Q1: Is this secure?

- It is *perfectly secure* (as long as keys are used only once)

### Q2: Why is this not used to encrypt everything?

- Keys must have *the same size* as the messages
  - To send a 2 Gb file, I must use a 2 Gb key!
- How can we pre-share and store such huge keys?
- But it is used everywhere in cryptography as a building block

## Kerckhoffs's Principle

- Long ago, it was common for encryption systems to be secret
- The idea is: the less people know, the harder it is to attack
- Also known as *Security through obscurity*
- We now know that **this is a bad idea**

## Kerckhoffs's Principle

- Long ago, it was common for encryption systems to be secret
- The idea is: the less people know, the harder it is to attack
- Also known as *Security through obscurity*
- We now know that **this is a bad idea**

### Kerckhoffs's Principle

- All details of a cryptosystem's operation are public
- The only secret is the key
- Why? Public knowledge promotes scrutiny
  - Designs of systems we will study are all public knowledge
  - Cryptographic schemes can be analyzed by everyone
  - Real-world security built on top of open standards
  - Methodology that revolutionized the way we approach security



## Never Use Your Own Crypto

A point we will hammer home.

- It is very easy to make mistakes
- It is very hard to find mistakes

## Never Use Your Own Crypto

A point we will hammer home.

- It is very easy to make mistakes
- It is very hard to find mistakes

### A plethora of bad apples

- Cryptography can be poorly designed
  - 2G mobile phone standard using A5/2 encryption
  - Broken after 10 years

## Never Use Your Own Crypto

A point we will hammer home.

- It is very easy to make mistakes
- It is very hard to find mistakes

### A plethora of bad apples

- Cryptography can be poorly designed
  - 2G mobile phone standard using A5/2 encryption
  - Broken after 10 years
- Cryptography may not match the application
  - TLS servers try to decrypt incoming traffic
  - What if errors are observed on the network?
  - Security definition does not capture this
  - Padding-oracle attacks are an example of this mismatch

## Never Use Your Own Crypto

A point we will hammer home.

- It is very easy to make mistakes
- It is very hard to find mistakes

### A plethora of bad apples

- Cryptography can be poorly designed
  - 2G mobile phone standard using A5/2 encryption
  - Broken after 10 years
- Cryptography may not match the application
  - TLS servers try to decrypt incoming traffic
  - What if errors are observed on the network?
  - Security definition does not capture this
  - Padding-oracle attacks are an example of this mismatch
- Cryptography can be poorly implemented
  - Timing attacks used to break theoretically secure crypto
  - Implementation errors can leak secret keys (e.g. heartbleed)

## Key Takeaways

- Encryption is a combination of two main algorithms:
  - Encryption takes plaintexts and produces ciphertexts
  - Decryption takes ciphertexts and produces plaintexts

## Key Takeaways

- Encryption is a combination of two main algorithms:
  - Encryption takes plaintexts and produces ciphertexts
  - Decryption takes ciphertexts and produces plaintexts
- Classical ciphers fail because of multiple reasons:
  - Small key space - key is easy to guess
  - Ciphertexts reveal patterns in the original message
  - One-time pad, on the other hand, is *perfectly secure*
  - But it is very *inefficient* to use in practice

## Key Takeaways

- Encryption is a combination of two main algorithms:
  - Encryption takes plaintexts and produces ciphertexts
  - Decryption takes ciphertexts and produces plaintexts
  
- Classical ciphers fail because of multiple reasons:
  - Small key space - key is easy to guess
  - Ciphertexts reveal patterns in the original message
  - One-time pad, on the other hand, is *perfectly secure*
  - But it is very *inefficient* to use in practice
  
- Kerckhoff's Principle
  - Algorithms should always be open-source
  - Security comes from the strength of the key
  - Many systems (today) rely on closed-source crypto



## Key Takeaways



- Encryption is a combination of two main algorithms:
  - Encryption takes plaintexts and produces ciphertexts
  - Decryption takes ciphertexts and produces plaintexts
- Classical ciphers fail because of multiple reasons:
  - Small key space - key is easy to guess
  - Ciphertexts reveal patterns in the original message
  - One-time pad, on the other hand, is *perfectly secure*
  - But it is very *inefficient* to use in practice
- Kerckhoff's Principle
  - Algorithms should always be open-source
  - Security comes from the strength of the key
  - Many systems (today) rely on closed-source crypto
- Do not write your own crypto!
  - It's easy to f-up
  - Testing correctness and security is very nuanced



## Defining Block Ciphers

A block cipher is defined by two deterministic algorithms

Encrypt:  $E(k, p)$

- Takes a key  $k \in \{0, 1\}^\lambda$
- Takes a plaintext block  $p \in \{0, 1\}^B$
- Outputs a ciphertext block  $c \in \{0, 1\}^B$

## Defining Block Ciphers

A block cipher is defined by two deterministic algorithms

Encrypt:  $E(k, p)$

- Takes a key  $k \in \{0, 1\}^\lambda$
- Takes a plaintext block  $p \in \{0, 1\}^B$
- Outputs a ciphertext block  $c \in \{0, 1\}^B$

Decrypt:  $D(k, c)$

- Takes a key  $k \in \{0, 1\}^\lambda$
- Takes a ciphertext block  $c \in \{0, 1\}^B$
- Outputs a plaintext block  $p \in \{0, 1\}^B$

A block cipher is **invertible**:  $k$  defines a **permutation**

# Advanced Encryption Standard (AES)

## AES was standardized in 2000

- DES was still standard (56-bit keys)
- 3DES was a common solution for short keys (112-bit security)
- 3DES: use DES 3 times with 3 independent keys
- 3DES chains  $E(k_1, D(k_2, E(k_3, p)))$

# Advanced Encryption Standard (AES)

## AES was standardized in 2000

- DES was still standard (56-bit keys)
- 3DES was a common solution for short keys (112-bit security)
- 3DES: use DES 3 times with 3 independent keys
- 3DES chains  $E(k_1, D(k_2, E(k_3, p)))$

AES is now the most used block cipher, by far


- Available in mainstream CPUs as HW implementation

Selected as a result of a competition

- 1997-2000 public competition run by NIST
- This process has since become the norm
- Criteria: performance and resistance to cryptanalysis

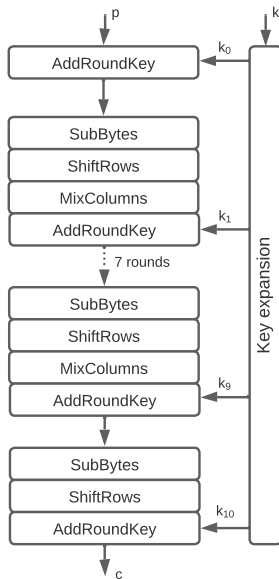
## Internals of AES

- Block size 128-bits and varying key size (128, 192, 256)-bits
- Keeps a 128-bit internal state: 4 x 4 array of 16-bits
- State is transformed using a substitution-permutation network




Substitutions/permutations have an algebraic description

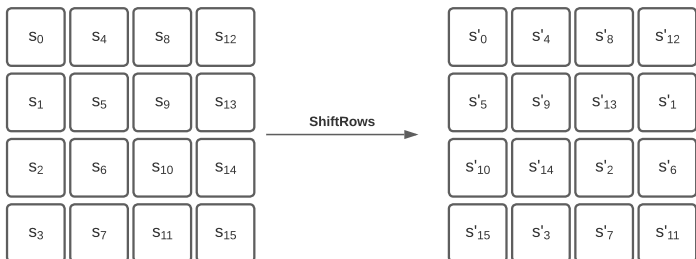
## Internals of AES - High Level View



## Internals of AES - SubBytes

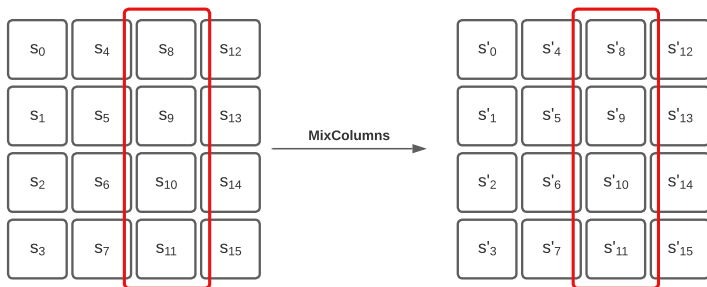


## Internals of AES - ShiftRows





## Internals of AES - MixColumns



## Using Block Ciphers Directly

Recall our secure block cipher building block:

Encrypt:  $E(k, p)$

- Takes a key  $k \in \{0, 1\}^\lambda$
- Takes a plaintext block  $p \in \{0, 1\}^B$
- Outputs a ciphertext block  $c \in \{0, 1\}^B$

Decrypt:  $D(k, c)$

- Takes a key  $k \in \{0, 1\}^\lambda$
- Takes a ciphertext block  $c \in \{0, 1\}^B$
- Outputs a plaintext block  $p \in \{0, 1\}^B$

**Q: What issues arise when using this to encrypt messages?**

## Modes of Operation

Modern cryptography clearly defines these concepts

- Block-ciphers are a **primitive**
- On their own, they're not very useful
- There are **insecure** ways to encrypt with a block cipher
- Encryption schemes have their own security definitions
- Encryption schemes built from block ciphers
- We prove encryption secure assuming a block cipher is secure

# Insecure Encryption from Secure Block Ciphers

## Electronic-Code-Book Mode (ECB)

- Break message into plaintext blocks  $p_0, \dots, p_n$
- Last block may need padding
  - That's a can of worms in and of itself
  - More on that later
- Independently encrypt each block  $c_i \leftarrow E(k, p_i)$

# Insecure Encryption from Secure Block Ciphers

## Electronic-Code-Book Mode (ECB)

- Break message into plaintext blocks  $p_0, \dots, p_n$
- Last block may need padding
  - That's a can of worms in and of itself
  - More on that later
- Independently encrypt each block  $c_i \leftarrow E(k, p_i)$
- **Q: Why is this insecure?**

# Insecure Encryption from Secure Block Ciphers

## Electronic-Code-Book Mode (ECB)

- Break message into plaintext blocks  $p_0, \dots, p_n$
- Last block may need padding
  - That's a can of worms in and of itself
  - More on that later
- Independently encrypt each block  $c_i \leftarrow E(k, p_i)$
- **Q: Why is this insecure?**


ECB is broken because you can see the penguin!

# Insecure Encryption from Secure Block Ciphers

## Electronic-Code-Book Mode (ECB)

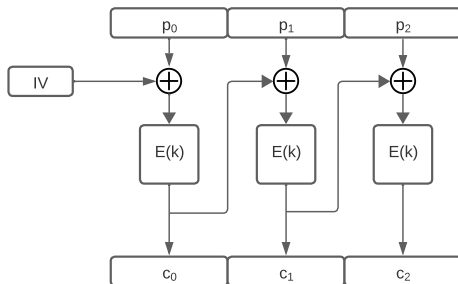
- Break message into plaintext blocks  $p_0, \dots, p_n$
- Last block may need padding
  - That's a can of worms in and of itself
  - More on that later
- Independently encrypt each block  $c_i \leftarrow E(k, p_i)$
- **Q: Why is this insecure?**

ECB is broken because you can see the penguin!



## Cipher Block Chaining

Engineers designed a secure encryption scheme before security proofs were well understood



- Main difference to ECB is the Initialization Vector (IV)
- Blocks depend on each other



## CBC: Padding

There are several padding methods

- Some schemes require message size as multiple of block size
- Padding schemes re-encode message so that is true
- To avoid ambiguity: **padding is always added**

## CBC: Padding

There are several padding methods

- Some schemes require message size as multiple of block size
- Padding schemes re-encode message so that is true
- To avoid ambiguity: **padding is always added**

The most common padding scheme is specified in PKCS#7:

- Let  $k > |M|$  be the next multiple of  $B$  (in bytes)
- Add  $k - |M|$  bytes with value  $k - |M|$
- The last byte always reveals how much padding was added
  - 0x01 means 1 byte of padding with that value
  - 0x03 means 3 bytes of padding with that value

## CBC: Padding

There are several padding methods

- Some schemes require message size as multiple of block size
- Padding schemes re-encode message so that is true
- To avoid ambiguity: **padding is always added**

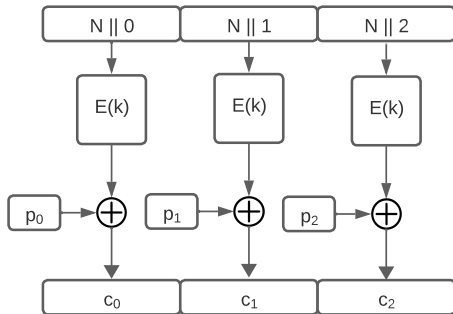
The most common padding scheme is specified in PKCS#7:

- Let  $k > |M|$  be the next multiple of  $B$  (in bytes)
- Add  $k - |M|$  bytes with value  $k - |M|$
- The last byte always reveals how much padding was added
  - 0x01 means 1 byte of padding with that value
  - 0x03 means 3 bytes of padding with that value

**Q: What is the minimum and maximum of added padding?**

## Counter Block Mode

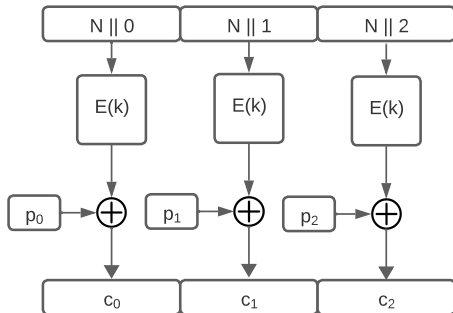
Often Counter Block Mode (CTR) is used in Nonce-based form



- $N$  must be unique, but not necessarily random
- Encryption becomes stateful

## Counter Block Mode

Often Counter Block Mode (CTR) is used in Nonce-based form



- $N$  must be unique, but not necessarily random
- Encryption becomes stateful
- **Q: How can this be faster than CBC?**

## Advantages of CTR

Counter mode is very efficient

- Key stream can be pre-processed
  - Block cipher not applied to the message!
- Any part of the data can be accessed efficiently
- This includes read/write access
- Decryption/encryption can be parallelized

As such, many modern protocols rely on CTR mode

## Key Takeaways

- AES selected via public competition
  - ... as all modern ciphers are

## Key Takeaways

- AES selected via public competition
  - ... as all modern ciphers are
- SubBytes; ShiftRows; MixColumns; AddRoundKey



## Key Takeaways

- AES selected via public competition
  - ... as all modern ciphers are
- SubBytes; ShiftRows; MixColumns; AddRoundKey
- Currently the *de facto* standard for block ciphers

## Key Takeaways

- AES selected via public competition
  - ... as all modern ciphers are
- SubBytes; ShiftRows; MixColumns; AddRoundKey
- Currently the *de facto* standard for block ciphers
- Block ciphers by themselves are **insecure**

## Key Takeaways

- AES selected via public competition
  - ... as all modern ciphers are
- SubBytes; ShiftRows; MixColumns; AddRoundKey
- Currently the *de facto* standard for block ciphers
- Block ciphers by themselves are **insecure**
- So we rely on modes of encryption: ECB, CBC, CTR

## On Keys

- All of this uses and generates keys
- How is this done?

## On Keys

- All of this uses and generates keys
- How is this done?

### For Symmetric Crypto

- Generated uniformly at random
- Derived using a Key Derivation Function
  - From a password or low entropy secret
  - From a high-entropy master key from key exchange protocol

## On Keys

- All of this uses and generates keys
- How is this done?

### For Symmetric Crypto

- Generated uniformly at random
- Derived using a Key Derivation Function
  - From a password or low entropy secret
  - From a high-entropy master key from key exchange protocol

### For Asymmetric Crypto

- Key generation algorithm  $\rightarrow$  key pair
- Private key holder generates both keys; publishes public key
- Asymmetric keys are typically much larger
  - RSA keys take roughly 4096-bits for 128-bit security
  - Elliptic-curve keys take roughly 400-bits for 128-bit security

## Storage and Generation

Keys are often *the most sensitive material* a secure system holds

## Storage and Generation

Keys are often *the most sensitive material* a secure system holds

Ideally, in an external secure hardware

- Hardware Security Module (HSM)
- Smartcard or similar cryptographic token



## Storage and Generation

Keys are often *the most sensitive material* a secure system holds

Ideally, in an external secure hardware

- Hardware Security Module (HSM)
- Smartcard or similar cryptographic token

### Key wrapping

- Long-term keys are often *wrapped* before storage
- To encrypt with another key
- Password-based encryption (low security)
- Wrap with HW-protected master key (standard security)
- Master key stored in trusted hardware (high security)

## To Be Random

**Q1: Which of these numbers are random?**

1. 00000000
2. 10101010
3. 00100100
4. 10011101

## To Be Random

### Q1: Which of these numbers are random?

1. 00000000 - Not random!
2. 10101010 - Not random (pattern)
3. 00100100 - Maybe not random?
4. 10011101 - Seems random...

## To Be Random

### Q1: Which of these numbers are random?

1. 00000000 - Not random!
2. 10101010 - Not random (pattern)
3. 00100100 - Maybe not random?
4. 10011101 - Seems random...

Randomness is not a property of a bit string, but rather:

- The bit generation process
- The bit string sampling procedure

## To Be Random

### Q1: Which of these numbers are random?

1. 00000000 - Not random!
2. 10101010 - Not random (pattern)
3. 00100100 - Maybe not random?
4. 10011101 - Seems random...

Randomness is not a property of a bit string, but rather:

- The bit generation process
- The bit string sampling procedure

**Q2: Which of these numbers will more likely appear in a fair randomness generator?**

## Randomness Distributions

Randomized processes described using *randomness distributions*.

We start with the **uniform distribution** over a finite field  $S$ .

A process  $U$  samples from the uniform distribution if

$$\forall s^* \in S, \Pr[s = s^* : s \leftarrow U] = \frac{1}{|S|}$$

## Randomness Distributions

Randomized processes described using *randomness distributions*.

We start with the **uniform distribution** over a finite field  $S$ .

A process  $U$  samples from the uniform distribution if

$$\forall s^* \in S, \Pr[s = s^* : s \leftarrow U] = \frac{1}{|S|}$$

**Q1: If we roll a fair dice, what is the probability of getting 1?**

## Randomness Distributions

Randomized processes described using *randomness distributions*.

We start with the **uniform distribution** over a finite field  $S$ .

A process  $U$  samples from the uniform distribution if

$$\forall s^* \in S, \Pr[s = s^* : s \leftarrow U] = \frac{1}{|S|}$$

**Q1: If we roll a fair dice, what is the probability of getting 1?**

$$\frac{1}{6} \approx 0.1667$$



## Randomness Distributions

Randomized processes described using *randomness distributions*.

We start with the **uniform distribution** over a finite field  $S$ .

A process  $U$  samples from the uniform distribution if

$$\forall s^* \in S, \Pr[s = s^* : s \leftarrow U] = \frac{1}{|S|}$$

**Q1: If we roll a fair dice, what is the probability of getting 1?**

$$\frac{1}{6} \approx 0.1667$$

**Q2: If we do a fair sampling of a byte, what is the probability of getting 00000000 or 10011101?**

## Randomness Distributions

Randomized processes described using *randomness distributions*.

We start with the **uniform distribution** over a finite field  $S$ .

A process  $U$  samples from the uniform distribution if

$$\forall s^* \in S, \Pr[s = s^* : s \leftarrow U] = \frac{1}{|S|}$$

**Q1: If we roll a fair dice, what is the probability of getting 1?**

$$\frac{1}{6} \approx 0.1667$$

**Q2: If we do a fair sampling of a byte, what is the probability of getting 00000000 or 10011101?**

$$\text{Both are } \frac{1}{2^8} \approx 0.0078$$

## Caution: statistical tests are not sufficient

- **Q: What type of tests can we do over “random” inputs?**

## Caution: statistical tests are not sufficient

- **Q: What type of tests can we do over “random” inputs?**
  - Count number of 1s and 0s
  - Check distribution of 8-bit words
  - Look for patterns
  - ...

### Irrelevant for Security

- Possible to pass statistical tests
- Totally insecure for cryptographic purposes

## Linux systems

- PRG is accessible at */dev/urandom*
  - In \*nix-style, PRG is mapped to a file
  - Careful to make sure system calls are successful!

## Linux systems

- PRG is accessible at */dev/urandom*
  - In \*nix-style, PRG is mapped to a file
  - Careful to make sure system calls are successful!

Link to code from **LibreSSL**

In some variants, there is a blocking */dev/random* based on an entropy simulator

- Check if there is “sufficient entropy”
- Blocks otherwise
- Current consensus indicates that, for most applications, this is not useful (see **this link** for more information)

# Concrete Numbers

## Some numbers for scale

- Not easy to perceive very very large numbers
- The estimated age of the universe in nanosecs is around  $2^{88}$
- The number of atoms in the universe is roughly  $2^{256}$

# Concrete Numbers

## Some numbers for scale

- Not easy to perceive very very large numbers
- The estimated age of the universe in nanosecs is around  $2^{88}$
- The number of atoms in the universe is roughly  $2^{256}$

## A common security parameter

- A common size for keys is 128 bits
- Consider the following events
  - Winning a lottery with 9 million participants (all of Portugal)
  - Guessing a  $2^{128}$  size key at the first try



# Concrete Numbers

## Some numbers for scale

- Not easy to perceive very very large numbers
- The estimated age of the universe in nanosecs is around  $2^{88}$
- The number of atoms in the universe is roughly  $2^{256}$

## A common security parameter

- A common size for keys is 128 bits
- Consider the following events
  - Winning a lottery with 9 million participants (all of Portugal)
  - Guessing a  $2^{128}$  size key at the first try

**Q1: Which event is more likely?**

# Concrete Numbers

## Some numbers for scale

- Not easy to perceive very very large numbers
- The estimated age of the universe in nanosecs is around  $2^{88}$
- The number of atoms in the universe is roughly  $2^{256}$

## A common security parameter

- A common size for keys is 128 bits
- Consider the following events
  - Winning a lottery with 9 million participants (all of Portugal)
  - Guessing a  $2^{128}$  size key at the first try

**Q1: Which event is more likely?**

**Q2: By how much?**

## Quantifying Security

Lower bound on the work required for a successful attack

Number of steps of the best attack

- $n$ -bits security
- Best attack to break the scheme requires  $2^n$  steps
- $n$ -bit keys cannot ever give more than  $n$ -bit security

## Quantifying Security

Lower bound on the work required for a successful attack

Number of steps of the best attack

- $n$ -bits security
- Best attack to break the scheme requires  $2^n$  steps
- $n$ -bit keys cannot ever give more than  $n$ -bit security
  - **Q1: Why?**

## Quantifying Security

Lower bound on the work required for a successful attack

### Number of steps of the best attack

- $n$ -bits security
- Best attack to break the scheme requires  $2^n$  steps
- $n$ -bit keys cannot ever give more than  $n$ -bit security
  - **Q1: Why?**
- Brute-force attack allows finding the correct key
- $l$ -bit keys could lead to  $n$ -bit security s.t.  $n \ll t$

## Quantifying Security

Lower bound on the work required for a successful attack

### Number of steps of the best attack

- $n$ -bits security
- Best attack to break the scheme requires  $2^n$  steps
- $n$ -bit keys cannot ever give more than  $n$ -bit security
  - **Q1: Why?**
- Brute-force attack allows finding the correct key
- $l$ -bit keys could lead to  $n$ -bit security s.t.  $n \ll t$ 
  - **Q2: When?**

## Quantifying Security

Lower bound on the work required for a successful attack

### Number of steps of the best attack

- $n$ -bits security
- Best attack to break the scheme requires  $2^n$  steps
- $n$ -bit keys cannot ever give more than  $n$ -bit security
  - **Q1: Why?**
- Brute-force attack allows finding the correct key
- $l$ -bit keys could lead to  $n$ -bit security s.t.  $n \ll t$ 
  - **Q2: When?**
  - Best attack is more efficient than brute-force
  - Common in asymmetric cryptography
  - Keys must follow specific structures, not random bit strings
- Quantifying using  $n$ -bit security permits comparing schemes

## Good Security Values for Real-world Crypto

The  $2^{128}$  rule of thumb

- Designs for which best attacks require roughly  $2^{128}$  attempts



# Good Security Values for Real-world Crypto

## The $2^{128}$ rule of thumb

- Designs for which best attacks require roughly  $2^{128}$  attempts

## For how long do we need security to hold?

- Moore's law: computational power doubles every 2 years
- $n + 1$  bit security every 2 years
- This no longer seems to be true, but...
- Maybe we will have quantum computers soon

Long-term security:  $\approx$  256-bit keys

Short-term security:  $\approx$  80-bit keys may be OK

## Key Takeaways

- Randomness is an important and challenging topic
  - Necessary to ensure strength of cryptographic keys
  - which (recall Kerckhoff's principle) is where security lies
  - It is a property of the generator, not the values

## Key Takeaways

- Randomness is an important and challenging topic
  - Necessary to ensure strength of cryptographic keys
  - which (recall Kerckhoff's principle) is where security lies
  - It is a property of the generator, not the values
- For real-world systems
  - Generators use entropy to gather randomness
  - Hard-to-predict events
  - Upon setup, entropy is low, so *be careful!*

## Key Takeaways

- Randomness is an important and challenging topic
  - Necessary to ensure strength of cryptographic keys
  - which (recall Kerckhoff's principle) is where security lies
  - It is a property of the generator, not the values
- For real-world systems
  - Generators use entropy to gather randomness
  - Hard-to-predict events
  - Upon setup, entropy is low, so *be careful!*
- Concrete security
  - Security relates to the size of the key
  - Maximum security is the key size (recall one-time-pad!)
  - $2^{128}$  bit security is often a good number
  - Long-term: 256-bit keys; short-term 80-bit sometimes suffices

# Computer Security Foundations

## Week 8: Symmetric Encryption

Bernardo Portela

L.EIC - 24