Computer Security Foundations Week 9: MACs and Authenticated Encryption

Bernardo Portela

L.EIC - 24

What is a Hash Function?

Hash functions are everywhere

- Key derivation
- Digest for authentication
- Randomness extraction
- Password protection
- Proofs of work

What is a Hash Function?

Hash functions are everywhere

Key derivation

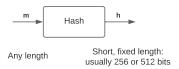
Hash Functions

•00000

- Digest for authentication
- Randomness extraction
- Password protection
- Proofs of work

Not only in crypto:

- Indexing in version management
- Deduplication in cloud storage systems
- File integrity in intrusion detection



000000

Secure Cryptographic Hash Functions

Efficient algorithms with nice properties

- Unpredictable outputs
- Hard to find pre-images
- Hard to find collisions

Secure Cryptographic Hash Functions

Efficient algorithms with nice properties

- Unpredictable outputs
- Hard to find pre-images
- Hard to find collisions

Hash functions are validated heuristically

- Similar to process for AES
- International competition for select designs
- Competitors are scrutinized wrt security and performance
- Several rounds, so more eyes on small number of proposals
- Most recent one: SHA-3

Finding Collisions

Collisions can be found with work $\sqrt{2^n}$, much better than 2^n !

Methodology

- Compute values like the brute-force attack
- Store them in a data structure indexed by image value
- Each new image value is searched in data structure
- Repeat until a collision is found

Finding Collisions

Collisions can be found with work $\sqrt{2^n}$, much better than 2^n !

Methodology

Hash Functions

- Compute values like the brute-force attack
- Store them in a data structure indexed by image value
- Each new image value is searched in data structure
- Repeat until a collision is found

How many operations?

• After *n* values, we checked n*(n-1)/2 pairs **Q: why?**

Finding Collisions

Collisions can be found with work $\sqrt{2^n}$, much better than 2^n !

Methodology

Hash Functions

- Compute values like the brute-force attack
- Store them in a data structure indexed by image value
- Each new image value is searched in data structure
- Repeat until a collision is found

How many operations?

- After *n* values, we checked n*(n-1)/2 pairs **Q: why?**
- Checking 2^n pairs takes roughly $\sqrt{2^n}$ values
- Overall complexity is that of finding the pre-image of a hash with n/2 bits of output (only half of the range)

The birthday paradox (not very paradoxical, just counterintuitive)

Building Hash Functions

Two main approaches that use iterative processes

 Merkle-Damgård construction: Used for MD4, MD5, SHA-1, SHA-256, SHA-512. Relies on a m + n-to-n bits compression function to construct a hash function of output length n for arbitrary input lengths

Building Hash Functions

Two main approaches that use iterative processes

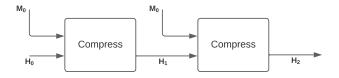
- Merkle-Damgård construction: Used for MD4, MD5, SHA-1, SHA-256, SHA-512. Relies on a m + n-to-n bits compression function to construct a hash function of output length n for arbitrary input lengths
- Sponge construction: Used for SHA-3, uses a *I*-bit permutation to construct a hash function for arbitrary input and output lengths

000000

Merkle-Damgård Construction

All prominent hash functions from 80s-2000s.

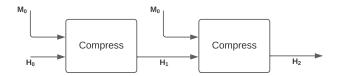
- H_0 is the initial value: constant and **public**
- M is broken into blocks of size m, M_1 , M_2 , . . .



Merkle-Damgård Construction

All prominent hash functions from 80s-2000s.

- *H*₀ is the initial value: constant and **public**
- M is broken into blocks of size m, M_1 , M_2 , . . .



- SHA-256: block size 512, output size 256 bits
- SHA-512: block size 1024, output size 512 bits
- What if messages are not of the same size as the block?

Sponge Construction

Absorb

- Fixed initial value h_0 , gradually accumulate message into state
- Message broken in blocks of size r (rate)
- Block XOR'ed into state

7 / 41

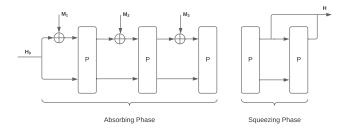
Sponge Construction

Absorb

- Fixed initial value h_0 , gradually accumulate message into state
- Message broken in blocks of size r (rate)
- Block XOR'ed into state

Squeeze

- Dual process iteratively constructs output
- Output constructed block by block



MD5

- Broken! 128-bit output
- Most popular hash function until broken in 2005
- These days, it takes seconds to find collisions
- The SHA function family (next) uses a similar design

Secure Hash Function (SHA)

Standardized by NIST in the US. International *de facto* standard SHA-0 published in 93', replaced with SHA-1 in 95'

- Both with 160-bit outputs
- Vulnerability not public at the time
- Later discovered collision attack in $2^{60} << 2^{80}$ operations
- More recent attacks reduced it to 2³³

Standardized by NIST in the US. International *de facto* standard SHA-0 published in 93', replaced with SHA-1 in 95'

- Both with 160-bit outputs
- Vulnerability not public at the time
- Later discovered collision attack in $2^{60} << 2^{80}$ operations
- More recent attacks reduced it to 2³³

SHA-1 remained unbroken until quite recently -(2017)

Most applications currently use SHA-2 (256 or 512 bits)

• Same design principles; larger parameters

Future applications adopting SHA-3 evolve to the Sponge

• Flexible output size is very useful!

SHA-1 Internals

- Merkle-Damgård, with Davis-Meyer compression function
- Block cipher used in compression function called SHACAL
 - Block cipher with 160-bit block sizes!

SHA-1 Internals

- Merkle-Damgård, with Davis-Meyer compression function
- Block cipher used in compression function called SHACAL
 - Block cipher with 160-bit block sizes!
- Message blocks are 512-bits, hashes are 160-bits long
- Davis-Meyer addition (not XOR): five 32-bit additions
- Insecure! Expected collisions in 2⁶³ ops in 2015, found in 2017

SHA-1 Internals

- Merkle-Damgård, with Davis-Meyer compression function
- Block cipher used in compression function called SHACAL
 - Block cipher with 160-bit block sizes!
- Message blocks are 512-bits, hashes are 160-bits long
- Davis-Meyer addition (not XOR): five 32-bit additions
- Insecure! Expected collisions in 2⁶³ ops in 2015, found in 2017

```
SHA1-blockcipher(a, b, c, d, e, M) {
    W = expand(M);
    for i = 0 to 79 { // K are constants
        new = (a <<< 5) + f(i, b, c, d) + e + K[i] + W[i]
        (a, b, c, d, e) = (new, a, b >>> 2, c, d)
    }
    return (a, b, c, d, e)
}
```

SHA-2 Family

- Family of 4 hash functions
 - SHA-224; 256; 384; 512

SHA-2 Family

- Family of 4 hash functions
 - SHA-224; 256; 384; 512
- Three digit identifier defines the output length
- Increased parameters and improved internal block ciphers
- SHA-224 and 256 still use 512 bit blocks (64 rounds)
 - SHA-224 is exactly the same as SHA-256, but has different IV and truncated output
 - SHA-384 and SHA-512 are similarly related
- SHA-512 compression function very similar, but has 80 rounds

SHA-2 Family

- Family of 4 hash functions
 - SHA-224; 256; 384; 512
- Three digit identifier defines the output length
- Increased parameters and improved internal block ciphers
- SHA-224 and 256 still use 512 bit blocks (64 rounds)
 - SHA-224 is exactly the same as SHA-256, but has different IV and truncated output
 - SHA-384 and SHA-512 are similarly related
- SHA-512 compression function very similar, but has 80 rounds

No non-generic attacks exist on these hash functions

- Still SHA-3 was (prudently) developed with different design
- Also has the benefit of varying sized outputs
- Good to generate keys!

SHA-3

- Keccack selected in 2009
- 3-year NIST SHA-3 competition
- Competition called for new design, if SHA-2 gets attacked

SHA-3

- Keccack selected in 2009
- 3-year NIST SHA-3 competition
- Competition called for new design, if SHA-2 gets attacked

Keccack is very different and very flexible

- Sponge based with 1600-bits permutation (in SHA-3)
- Blocks can be 1152, 1088, 832 or 576 bits
- Corresponding to 224, 256, 384 or 512 bit outputs
- As a bonus we get the SHAKE functions
 - SHAKE128 and SHAKE256
 - eXtendable Output Functions (XOFs)
 - You can specify output length



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x
- For output of 2^n , collision can be found in $\approx \sqrt{2^n} = 2^{\frac{n}{2}}$
 - So for 2¹²⁸ resistance, output must be at least 2²⁵⁶



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x
- For output of 2^n , collision can be found in $\approx \sqrt{2^n} = 2^{\frac{n}{2}}$
 - So for 2¹²⁸ resistance, output must be at least 2²⁵⁶
- Two main constructions: MD and Sponge



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x
- For output of 2^n , collision can be found in $\approx \sqrt{2^n} = 2^{\frac{n}{2}}$
 - So for 2¹²⁸ resistance, output must be at least 2²⁵⁶
- Two main constructions: MD and Sponge
- Merkle-Damgård
 - Used in MD5; SHA1; SHA2



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x
- For output of 2^n , collision can be found in $\approx \sqrt{2^n} = 2^{\frac{n}{2}}$
 - So for 2¹²⁸ resistance, output must be at least 2²⁵⁶
- Two main constructions: MD and Sponge
- Merkle-Damgård
 - Used in MD5; SHA1; SHA2
- Sponge
 - Used in SHA3 and SHAKE



- Hash functions are one-way functions
 - From any sized inputs to fixed-size output
 - Easy to go from x to f(x)
 - Hard to go from f(x) to x
- For output of 2^n , collision can be found in $\approx \sqrt{2^n} = 2^{\frac{n}{2}}$
 - So for 2¹²⁸ resistance, output must be at least 2²⁵⁶
- Two main constructions: MD and Sponge
- Merkle-Damgård
 - Used in MD5; SHA1; SHA2
- Sponge
 - Used in SHA3 and SHAKE
- SHA-2 and SHA-3 currently the de facto standards

MACs as Keyed Hashes

Short Summaries of Potentially Large Messages

- Called a hash if everything is public
- Keyed hashes allows for conditional hash computation

Short Summaries of Potentially Large Messages

- Called a hash if everything is public
- Keyed hashes allows for conditional hash computation

Message Authentication Codes - MACs

- Symmetric Authentication $t \leftarrow MAC(k, m)$
- t guarantees that m was produced by someone that knows k
- Implies message *m* was not changed since its creation
- Digital signatures in the symmetric paradigm!

Message Authentication Codes

Typical use of MACs – SSH, IPSec, TLS

- Two parties was message authentication and integrity
- Some form of set-up/agreement to establish common key k
- Sender computes $t \leftarrow \mathsf{MAC}(k, m)$ and sends (m, t)
- Receiver gets (m, t), recomputes $t' \leftarrow MAC(k, m)$
- If $t \neq t'$, message is rejected!

Message Authentication Codes

Typical use of MACs – SSH, IPSec, TLS

- Two parties was message authentication and integrity
- Some form of set-up/agreement to establish common key k
- Sender computes $t \leftarrow \mathsf{MAC}(k, m)$ and sends (m, t)
- Receiver gets (m, t), recomputes $t' \leftarrow MAC(k, m)$
- If $t \neq t'$, message is rejected!

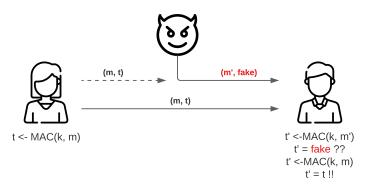
Acceptance means m was produced while knowing k

In this process, message is public!

MACs do not give confidentiality. They provide integrity

Its orthogonal to encryption. In real-world applications, we will need to combine these

Authentication and Message Integrity



- No possibility of computing t without k implies
- Adversary cannot change the message
- Adversary cannot conjure new messages

Some Context

MACs constructed from hash functions and block ciphers

Simplest construction: prefix key

$$MAC(k, m) = H(K||M)$$

Some Context

MACs constructed from hash functions and block ciphers

Simplest construction: prefix key

$$MAC(k, m) = H(K||M)$$

MD yields insecure MAC!

- Given (m, t), attacker outputs H(K||M||pad||M')
- This can be computed just from t' and m'
- Length extension attack

Some Context

MACs constructed from hash functions and block ciphers

Simplest construction: prefix key

$$MAC(k, m) = H(K||M)$$

MD yields insecure MAC!

- Given (m, t), attacker outputs H(K||M||pad||M')
- This can be computed just from t' and m'
- Length extension attack

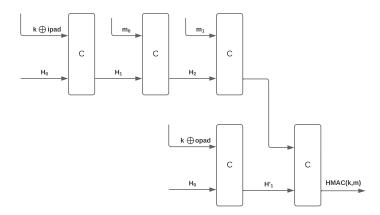
A consideration in SHA-3 construction

- Abandon MD construction
- Include explicit keyed hash

HMAC Construction

When instantiated with MD construction

- HMAC is simply $H((K \oplus opad)||H((k \oplus ipad)||m))$
- ipad and opad are constraints: align to block size



Building MACs from Block Ciphers

We have seen block ciphers \rightarrow hash functions \rightarrow MACs

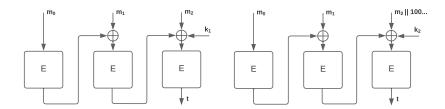
But there are also direct constructions: block ciphers \rightarrow MACs

CMAC

- Used in IPSec
- CMAC improves on CBC-MAC (which was broken!)
- Use CBC mode of operation
- Fix IV to all zero blocks
- Take the last ciphertext block as a tag

CMAC fixes CBC-MAC by processing last block differently

- All blocks except last are processed like CBC-MAC
- Keys k_1 and k_2 derived from k
 - $I \leftarrow E(k,0)$
 - $k_1 = (I << 1) \oplus (0 \times 00..0087 * LSB(I)))$
 - $k_2 = (k_1 << 1) \oplus (0 \times 00..0087 * LSB(k_1)))$



Universal Hash Functions

UHF are a Weak form of Hashing

- Don't need to be collision resistance
- Parametrised by a key UH(k, m)
- Guarantee that, for two fixed messages $m_0 \neq m_1$:

$$\Pr[\mathsf{UH}(k,m_0) = \mathsf{UH}(k,m_1)] \le \epsilon$$

ullet Considering random k and very small ϵ

Universal Hash Functions

UHF are a Weak form of Hashing

- Don't need to be collision resistance
- Parametrised by a key UH(k, m)
- Guarantee that, for two fixed messages $m_0 \neq m_1$:

$$\Pr[\mathsf{UH}(k,m_0) = \mathsf{UH}(k,m_1)] \le \epsilon$$

ullet Considering random k and very small ϵ

No other security requirements \rightarrow easy to construct

We can use a universal hash function as a MAC

Provided that we only authenticate one message!

Wegman-Carter Construction

How to circumvent this limitation?

- Use a PRF to strengthen the UH
- Converts a UH into a fully secure MAC
- AES can fill the PRF role!

Wegman-Carter Construction

How to circumvent this limitation?

- Use a PRF to strengthen the UH
- Converts a UH into a fully secure MAC
- AES can fill the PRF role!

Intuition: Encrypt Universal Hash Value

$$\mathsf{UH}(k_1,m) \oplus \mathsf{PRF}(k_2,n)$$

- The full MAC key is (k_1, k_2)
- *n* is a public value that must never repeat
 - A.k.a. a nonce
- This can be kept as a counter, or generated at random

Poly1305-AES: Wegman-Carter in Practice

- Initial proposal used AES as the Wegman-Carter PRF
- The universal hash function uses prime $2^{130} 5$

Poly1305
$$((k_1, k_2), m) = (m_1k + \ldots + m_nk^n \pmod{p}) + AES(k_2, n)$$

- Blocks are 128 bits and last block is padded with 100
- The final addition is performed modulo 2¹²⁸
- TLS recommends Poly1305 with ChaCha20, rather than AES



- Keyed hashing allows for message authentication
- For hash function h, one cannot produce h(k, x) w/o k
- Provides integrity; ciphers give confidentiality



- Keyed hashing allows for message authentication
- For hash function h, one cannot produce h(k, x) w/o k
- Provides <u>integrity</u>; ciphers give <u>confidentiality</u>
- Just add a key as prefix!
 - Problem! Length extension attacks



- Keyed hashing allows for message authentication
- For hash function h, one cannot produce h(k,x) w/o k
- Provides integrity; ciphers give confidentiality
- Just add a key as prefix!
 - Problem! Length extension attacks
- HMAC
 - Input padding and output padding, both using the key



- Keyed hashing allows for message authentication
- For hash function h, one cannot produce h(k,x) w/o k
- Provides integrity; ciphers give confidentiality
- Just add a key as prefix!
 - Problem! Length extension attacks
- HMAC
 - Input padding and output padding, both using the key
- CMAC
 - Do AES-CBC without IV; return the last block
 - With a twist to prevent prefix extension

🕏 Key Takeaways 🕏

- Keyed hashing allows for message authentication
- For hash function h, one cannot produce h(k,x) w/o k
- Provides integrity; ciphers give confidentiality
- Just add a key as prefix!
 - Problem! Length extension attacks
- HMAC
 - Input padding and output padding, both using the key
- CMAC
 - Do AES-CBC without IV; return the last block
 - With a twist to prevent prefix extension
- Wegman-Carter
 - Use a UHF for a unique message
 - XOR it with an encryption of a nonce
 - Used in AES-GCM (next!)

Why Authenticated Encryption?

Any secure channel in practice uses authenticated encryption

- Messages need to be confidential
- Messages need to be authentic
- Messages should not be repeated/omitted/removed

Why Authenticated Encryption?

Any secure channel in practice uses authenticated encryption

- Messages need to be confidential
- Messages need to be authentic
- Messages should not be repeated/omitted/removed

Encryption provides confidentiality

MACs provide authenticity

Why Authenticated Encryption?

Any secure channel in practice uses authenticated encryption

- Messages need to be confidential
- Messages need to be authentic
- Messages should not be repeated/omitted/removed

Encryption provides confidentiality

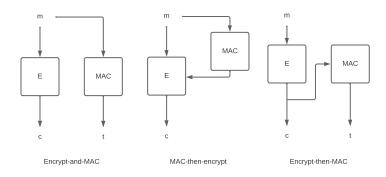
MACs provide authenticity

Authenticated public meta-information (e.g. sequence numbers) is used to solve the third point – Authenticated Encryption with Associated Data (AEAD)

Authenticated Encryption using MACs

The good, the bad and the ugly

- Encrypt-and-MAC: encryption and MAC of the message
- MAC-then-Encrypt: Encrypt message and its authentication
- Encrypt-then-MAC: Authenticate the message encryption



Encrypt-and-MAC

AE with $k = (k_1, k_2)$ done with parallel processing

- $c \leftarrow s E(k_1, m)$
- $t \leftarrow MAC(k_2, m)$
- Output: (*c*, *t*)

Encrypt-and-MAC

AE with $k = (k_1, k_2)$ done with parallel processing

- $c \leftarrow s E(k_1, m)$
- $t \leftarrow MAC(k_2, m)$
- Output: (*c*, *t*)

Problems - The bad

- Potentially malicious c decrypted before authentication
- MACs not designed to ensure confidentiality!
- Construction can be secure for some MACs...
 - Very easy to make mistakes
- Used in SSH: $MAC(k_2, m||n)$, where n is the sequence number

MAC-then-Encrypt

AE with $k = (k_1, k_2)$ done sequentially

- $t \leftarrow MAC(k_1, m)$
- $c \leftarrow_{\$} E(k_2, m||t)$
- Output: c

MAC-then-Encrypt

AE with $k = (k_1, k_2)$ done sequentially

- $t \leftarrow MAC(k_1, m)$
- $c \leftarrow s E(k_2, m||t)$
- Output: c

Problems – The ugly

- Potentially malicious c decrypted before authentication
- Used in TLS until version 1.3
- Painful story with padding oracle attacks
 - Issue arises from the decryption before authentication
 - Did the decryption fail because of the padding, or because of the MAC?
 - Theoretical attack found disregarded at first
 - Practical attack found a couple of years later: Lucky 13

Encrypt-then-MAC

AE with $k = (k_1, k_2)$ done sequentially

- $c \leftarrow s E(k_1, m)$
- $t \leftarrow MAC(k_2, c)$
- Output: c

Encrypt-then-MAC

AE with $k = (k_1, k_2)$ done sequentially

- $c \leftarrow s E(k_1, m)$
- $t \leftarrow MAC(k_2, c)$
- Output: c

Advantages - The good

- Ciphertext not decrypted unless it is authenticated
- Useful against DoS attacks
 - MAC verification typically very fast
- Preferred method, except in legacy systems

Authenticated Encryption from Scratch

Modern AEADs are not black-box compositions of Enc+MAC. Encryption/authentication layers visible in all constructions

Authenticated Encryption from Scratch

Modern AEADs are not black-box compositions of Enc+MAC. Encryption/authentication layers visible in all constructions

Optimized for Performance

- Encryption/decryption of blocks in parallel
- Throughput, streamability, memory requirements, input fixing:
 - Implementation cost, e.g., need block cipher inverse?
 - Can start transmitting before encryption is complete?
 - Can start decryption before ciphertext is fully received?
 - Can discard plaintext/ciphertext block immediately (online)?
 - Can metadata be given at any point? Or must be fixed initially?

Authenticated Encryption from Scratch

Modern AEADs are not black-box compositions of Enc+MAC. Encryption/authentication layers visible in all constructions

Optimized for Performance

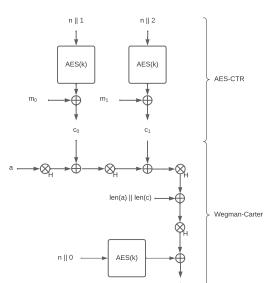
- Encryption/decryption of blocks in parallel
- Throughput, streamability, memory requirements, input fixing:
 - Implementation cost, e.g., need block cipher inverse?
 - Can start transmitting before encryption is complete?
 - Can start decryption before ciphertext is fully received?
 - Can discard plaintext/ciphertext block immediately (online)?
 - Can metadata be given at any point? Or must be fixed initially?

Streamability (a.k.a. online) \rightarrow low memory requirements

Very important for routers, TLS/https termination, etc.

Most widely used AEAD: IPSec, SSH, TLS

- GCM:
 - CTR plus auth layer
- Keys: 128 bits
- Nonce: 96 bits
- CTR starts at 1



Galois-Counter Mode: Authentication

Q: What type of AE is this?

Galois-Counter Mode: Authentication

Q: What type of AE is this? Encrypt-then-MAC!

MAC uses Wegman-Carter construction

- Universal Hash function is called GHASH
- AES used as PRF to hash the value on input $n \parallel 0$

Galois-Counter Mode: Authentication

Q: What type of AE is this? Encrypt-then-MAC!

MAC uses Wegman-Carter construction

- Universal Hash function is called GHASH
- AES used as PRF to hash the value on input $n \parallel 0$

$GHASH(hk, c_1, \ldots, c_n)$ defined as follows

- $hk \leftarrow AES(k,0)$
- Evaluate P(x) defined by (pad0(A), pad0(c), |A| || |C|) at hk
- Horner's formula: $P(x) = x * (x * (x * (...) + a_2) + a_1) + a_0$
- Algebraic magic:
 - Super efficient in hardware (special processor instructions)

Efficiency of Galois Counter Mode

Encryption layer inherits parallelism from CTR mode Authentication layer blocks if *a* is known only in the end Encryption layer inherits parallelism from CTR mode

Authentication layer blocks if a is known only in the end

- a known from the start \rightarrow GCM streamable
 - · Ciphertext blocks computed and authenticated on the fly
 - Authentication of previous block is accumulated while current block is being encrypted

Efficiency of Galois Counter Mode

Encryption layer inherits parallelism from CTR mode Authentication layer blocks if a is known only in the end

a known from the start \rightarrow GCM streamable

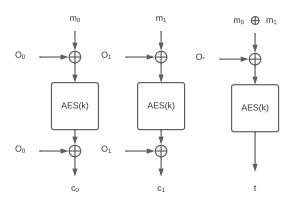
- Ciphertext blocks computed and authenticated on the fly
- Authentication of previous block is accumulated while current block is being encrypted

Could be faster...

- Authentication usually not computed fully in parallel
- HW implementations of auth layer slower than AES-CTR

(ChaCha20-Poly1305 not covered, but has similar structure)

Offset Codebook (OCB)



- Offset depends on the key and the nonce
- Incremented for each new block
- Last offset computed from last plaintext block processed

Using OCB as AEAD

- Very secure and efficient
- Implementations required licensing
- Patent-renewal fees intentionally not payed (see here)

Using OCB as AEAD

- Very secure and efficient
- Implementations required licensing
- Patent-renewal fees intentionally not payed (see here)

Similar to AES-GCM, OBC allows authenticated data a_0, a_1, \ldots, a_n Tag computed as

$$T = E(k, S \oplus O_*) \oplus E(k, a_0 \oplus O_0) \oplus E(k, a_1 \oplus O_1) \oplus \dots$$

where $S = m_0 \oplus m_1 \oplus \dots$

Offset values for AD different from those used to encrypt m_0, m_1, \ldots, m_n

Security and Efficiency of OCB

Nonce reuse

- On nonce reuse, the attacker can identify block duplicates
 - E.g. block 3 of message 1 is similar to block 5 of message 2
 - GCM allows detection of multiple blocks, but also XOR differences for blocks in the same position
- Repeated nonces can break the authenticity of OCB
 - An attacker can combine blocks from messages authenticated with OCB to create another authenticated message
 - ...but unlike GCM, it cannot extract the underlying key!!

Security and Efficiency of OCB

Nonce reuse

- On nonce reuse, the attacker can identify block duplicates
 - E.g. block 3 of message 1 is similar to block 5 of message 2
 - GCM allows detection of multiple blocks, but also XOR differences for blocks in the same position
- Repeated nonces can break the authenticity of OCB
 - An attacker can combine blocks from messages authenticated with OCB to create another authenticated message
 - ...but unlike GCM, it cannot extract the underlying key!!

Efficiency

- OCB and GCM make about as many calls to the block cipher
- GCM used to be 3x slower than OCB
 - AES and GHASH competed for CPU resources
- OCB requires encryption and decryption, contrary to GCM

Strengthening AES-GCM against nonce-reuse.

Generic composition of nonce-based encryption and a PRF

- $t \leftarrow PRF(k_1, a || p || n)$
- $c \leftarrow Enc(k_2, n = t, p)$
- Output (*c*, *t*)

Strengthening AES-GCM against nonce-reuse.

Generic composition of nonce-based encryption and a PRF

- $t \leftarrow PRF(k_1, a || p || n)$
- $c \leftarrow Enc(k_2, n = t, p)$
- Output (c, t)

Q1: What happens if?

• *n* repeats, but *a* or *p* changes?

Strengthening AES-GCM against nonce-reuse.

Generic composition of nonce-based encryption and a PRF

- $t \leftarrow PRF(k_1, a || p || n)$
- $c \leftarrow Enc(k_2, n = t, p)$
- Output (c, t)

Q1: What happens if?

- n repeats, but a or p changes?
- All of a, n and p repeat?

Strengthening AES-GCM against nonce-reuse.

Generic composition of nonce-based encryption and a PRF

- $t \leftarrow PRF(k_1, a || p || n)$
- $c \leftarrow Enc(k_2, n = t, p)$
- Output (*c*, *t*)

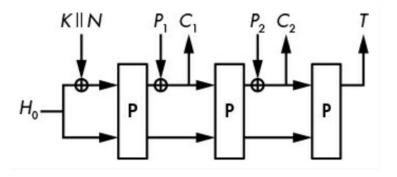
Q1: What happens if?

- n repeats, but a or p changes?
- All of a, n and p repeat?
- Tag used as encryption nonce is fresh w/ high probability
- Scheme is not streamable! Q2: Why?

Building AE from Permutations

Recall the sponge construction of SHA-3

Closely related construction - Duplex - gives an AEAD



- P is a fixed (unkeyed) permutation and h_0 is a public value
- Last block must be padded
- AEAD versions slightly more involved. Not covered

Permutation-based AEs

Resulting constructions are

- Fast
- Streamable

Permutation-based AEs

Resulting constructions are

- Fast
- Streamable

Interesting nonce reuse resilience:

- Unforgeability not affected
- Plaintexts compromised
 - First block
 - Subsequent block if there is a common prefix
- Plaintexts remain confidential after divergence



- Authenticated Encryption ensures <u>confidentiality</u> and <u>integrity</u>
- Encrypt-and-MAC
 - Exposed MAC. Not good!
- MAC-then-Encrypt
 - Possible decryption of malicious ciphertexts;
 - Not necessarily broken; has subtle issues
- Encrypt-then-MAC
 - The "safest" way to do it



- Authenticated Encryption ensures <u>confidentiality</u> and <u>integrity</u>
- Encrypt-and-MAC
 - Exposed MAC. Not good!
- MAC-then-Encrypt
 - Possible decryption of malicious ciphertexts;
 - Not necessarily broken; has subtle issues
- Encrypt-then-MAC
 - The "safest" way to do it
- General usage \rightarrow ambitious requirements



- Authenticated Encryption ensures confidentiality and integrity
- Encrypt-and-MAC
 - Exposed MAC. Not good!
- MAC-then-Encrypt
 - Possible decryption of malicious ciphertexts;
 - Not necessarily broken; has subtle issues
- Encrypt-then-MAC
 - The "safest" way to do it
- General usage \rightarrow ambitious requirements
- AES-GCM
 - AES-CTR followed by a round of Wegman-Carter
 - ChaCha20-Poly1305 similar

- Authenticated Encryption ensures confidentiality and integrity
- Encrypt-and-MAC
 - Exposed MAC. Not good!
- MAC-then-Encrypt
 - Possible decryption of malicious ciphertexts;
 - Not necessarily broken; has subtle issues
- Encrypt-then-MAC
 - The "safest" way to do it
- General usage \rightarrow ambitious requirements
- AES-GCM
 - AES-CTR followed by a round of Wegman-Carter
 - ChaCha20-Poly1305 similar
- OCB
 - Offset XOR'd with AES input/output



- Authenticated Encryption ensures confidentiality and integrity
- Encrypt-and-MAC
 - Exposed MAC. Not good!
- MAC-then-Encrypt
 - Possible decryption of malicious ciphertexts;
 - Not necessarily broken; has subtle issues
- Encrypt-then-MAC
 - The "safest" way to do it
- General usage → ambitious requirements
- AES-GCM
 - AES-CTR followed by a round of Wegman-Carter
 - ChaCha20-Poly1305 similar
- OCB
 - Offset XOR'd with AES input/output
- AEAD from permutations
 - Absorb phase from Sponge; gets blocks between permutations
 - Authentication tag recovered at the end

Computer Security Foundations Week 9: MACs and Authenticated Encryption

Bernardo Portela

L.EIC - 24